

HAIA-CAIPR

Cross AI Platform Review

A Governance Protocol for Human Orchestration of Parallel Multi-AI Execution

Basil C. Puglisi, MPA

Human-AI Collaboration Strategist

Author, Governing AI: When Capability Exceeds Control

basilpuglisi.com | github.com/basilpuglisi/HAIA

Revised April 2026

Executive Summary

HAIA-CAIPR (Cross AI Platform Review) is the governance protocol for human orchestration of parallel multi-AI execution within the HAIA framework. It sits between HAIA-RECCLIN and Checkpoint-Based Governance (CBG) in the governance stack. It emerged from eight months of daily operational practice across eleven AI platforms and was formally named and specified in March 2026, documented in Case Study 006.

RECCLIN tells each AI how to respond to the human, while CAIPR tells the human how to work across multiple AIs simultaneously. CBG tells the human when to decide with constitutional authority, and GOPEL will automate CAIPR's mechanics without performing cognitive work.

The distinction from RECCLIN is not incremental but architectural. RECCLIN governs one AI at a time, while CAIPR governs the human's orchestration of many AIs in parallel. The evidence problems that parallel orchestration creates, which include hallucination cross-validation, algorithmic narcissism, source-authority erosion, synthesizer failure, platform behavioral clustering, and convergent error risk, cannot be addressed at the single-AI level.

CAIPR was discovered through practice, not designed in advance. Four findings were predicted by theory, and twenty-nine were discovered through practice. That ratio is itself a finding. In this operational record, practice produced seven times more findings than initial design anticipated. Governance methodologies practiced at sufficient intensity and duration may surface discoveries that theoretical design cannot anticipate, though broader generalization requires replication outside the current record.

Evidence basis and scope. The operational record underlying this specification spans eight months of daily multi-platform practice across eleven AI platforms, applied to published work in multiple categories: two editions of a full-length book (Governing AI: When Capability Exceeds Control), a Congressional policy package (AI Provider Plurality v9, including the Verified AI Inference Standards Act and the Council for Humanity proposal), two academic submissions (SSRN 6195238 and SSRN 6583419), long-form white papers and research briefs, a blog archive exceeding one thousand posts, and the code development cycle for GOPEL including adversarial review by independent AI platforms. This specification is derived from that record. External replication in third-party operational contexts by practitioners outside the author's work remains a future milestone.

Part I. Why CAIPR Exists as a Methodology Beyond HAIA-RECCLIN

1.1 What RECCLIN Does and Does Not Govern

HAIA-RECCLIN governs the structured output format for any single human-AI interaction. Its seven roles (Researcher, Editor, Coder, Calculator, Liaison, Ideator, Navigator) define what the AI must declare when responding: its functional role, task understanding, evidence, conflicts, confidence, time-sensitivity of information, a Factics triad (Fact to Tactic to KPI), a primary recommendation, and the decision requiring human approval.

RECCLIN also performs a second function as important as the output format: it is a cognitive development system for the human governor. The structured output trains the human to evaluate rather than accept AI responses. Through repetition, the governor develops the capacity to read outputs critically, recognize unjustified confidence claims, surface hidden assumptions, and distinguish useful dissent from noise. This developed capacity is the prerequisite for everything that follows in the framework.

What RECCLIN does not govern: how the human works when multiple AIs run simultaneously. RECCLIN specifies how each platform should respond. It provides no guidance on how to dispatch to multiple platforms, compare their outputs, detect when one fabricates while others do not, manage platform behavioral differences, govern the AI synthesizing multiple outputs, or ensure the human's own input is recognized as constitutionally distinct when flowing through multi-platform workflows.

These are not edge cases. They are the defining operational challenges of multi-AI governance practice. Eight months of daily work across eleven platforms produced 33 documented findings. Twenty-eight of those 33 originated specifically in the parallel execution context RECCLIN was not designed to address.

1.2 The Failure Evidence That Required a New Layer

Hallucination Detection Failure (Finding 2.2). In single-platform workflows, there is no mechanism to detect fabrication. Cross-platform parallel execution exposes fabrications because other platforms do not corroborate inventions. Three documented incidents: wholesale document section fabrication, quote fabrication with self-review of those fabricated quotes, and content repetition fifteen times in a single output. Single-platform governance is structurally blind to this failure mode.

Source-Authority Erosion (The Loop That Ate the Governor). In a multi-AI workflow using a synthesizer, the human arbiter's own input was processed identically to AI platform output. The synthesizer treated human corrections as one more data source. The human arbiter began doubting their own corrections. This is reverse automation bias: the human doubting their own authority because the system failed to recognize it.

Algorithmic Narcissism (Finding 2.1). In a nine-platform and separately a three-platform structured review, every platform asked to identify which should synthesize claimed itself with supporting evidence. This is a structural property of AI systems asked to evaluate their own fitness. RECCLIN provides no mechanism to prevent a single-platform synthesizer chokepoint from forming by default.

1.3 Why CAIPR Is a Methodology, Not an Extension

An extension modifies or expands an existing system. A methodology governs a distinct operational domain with its own principles, failure modes, and required practices.

CAIPR governs a domain RECCLIN cannot reach: the space between multiple AI platforms and the human arbiter extracting disciplined intelligence from them simultaneously. That space has its own structural properties: convergence that is not confirmation, dissent that is not error, synthesizers that fail invisibly, platforms that behave systematically differently, and human authority that can vanish without detection if not classified at ingestion.

The operational analogy: RECCLIN is the grammar each witness speaks when testifying. CAIPR is the protocol governing how a deliberative body receives, compares, and evaluates multiple witnesses simultaneously. Both are necessary. Neither can substitute for the other.


Critically, CAIPR cannot correct at the constitutional level. It corrects at the output level by comparing responses across platforms. Training data, alignment tuning, and platform values are set before user interaction. Multi-platform comparison reveals where platform constitutions diverge. It cannot change them (Finding 2.14).

1.4 The Naming Decision

The acronym CAIPR is human-arbiter originated. No AI platform produced it.

When multiple platforms were asked to describe what this layer actually does, every platform produced the same functional description: it governs cross-platform review. The human arbiter synthesized that convergence with a direct naming principle. Call the layer what it does. Cross AI Platform Review. CAIPR.

The acronym has no known collision in the AI governance space, carries no surveillance connotations, claims no authority belonging to CBG, and is self-describing. Human-originated naming of a human-centric protocol is itself a governance signal.



Part II. Core Definition and Scope

2.1 Definition

HAIA-CAIPR (Cross AI Platform Review) is the governance protocol for human orchestration of parallel multi-AI execution within the HAIA framework. It specifies how a human governor dispatches identical prompts to multiple independent AI platforms, collects their RECCLIN-structured outputs, compares those outputs for convergence and divergence, detects hallucinations and fabrications through cross-validation, governs the AI synthesizing multiple platform outputs, and maintains constitutional source-authority discrimination throughout.

CAIPR operates under CBG constitutional authority, and every checkpoint in a CAIPR-governed session is a CBG checkpoint. CAIPR does not create new governance authority; it specifies how CBG authority is exercised in the multi-AI context.

Governance meta-requirement. Every emergent CAIPR protocol refinement must be reviewed against CBG constitutional principles before adoption as standard practice. The sequential intake protocol failure (Finding 2.33) documented a structural violation that occurred because an operationally-invented protocol was adopted without constitutional review. Protocol 6.8 closes that vulnerability.

2.2 The Framework Position

- **Facts** foundational evidentiary discipline (2012)
- **HAIA-RECCLIN** structured AI-response grammar (single-platform)
- **HAIA-CAIPR** human-orchestration grammar (multi-platform). This document.
- **CBG** constitutional checkpoint authority
- **HEQ and AIS** human cognitive growth measurement
- **HAIA-Agent** automated orchestration logistics
- **GOPEL** non-cognitive enforcement infrastructure

2.3 The Three Practitioner Operating Modes

These modes describe where a practitioner is in the adoption progression. They are distinct from the three CBG operating models, which describe checkpoint density and automation level.

Learning Mode. Role assignment to single platforms. Useful for onboarding practitioners building RECCLIN governance muscle before scaling to parallel execution.

Operational Mode. Parallel execution with identical prompts to all available platforms. Current standard practice. All 33 operational findings were produced in this mode.

Agent Mode. Role-mapped sub-agents in automated pipelines with GOPEL executing CAIPR mechanics. Specified but not yet deployed as production software.

2.4 Relationship to the Three CBG Operating Models


Model 1, Agent Responsible AI. One CBG checkpoint at final output. Appropriate for low-to-moderate risk tasks.

Model 2, Agent AI Governance. CBG checkpoint after each RECCLIN role. Higher governance density.

Model 3, Manual Human AI Governance. Human performs all CAIPR operations directly. No agent. Current standard practice because GOPEL is not yet deployed. All 33 operational findings were produced here.

The HAIA Framework Stack

Where CAIPR sits in the governance architecture



Part III. The Eight Core Operations

3.1 Parallel Dispatch

The same prompt is sent to multiple independent AI platforms simultaneously. Sequential dispatch allows early outputs to contaminate later prompting. Parallel dispatch preserves independence.

Platform count follows odd-number protocol: 3, 5, 7, 9, or 11 platforms. Minimum for any CAIPR session is 3 platforms.

Platform count matches stakes, not an efficiency ceiling. An early operating principle suggested diminishing returns past six or seven platforms. Two documented events overrode it. In one review, the tenth platform produced the most rigorous output in the corpus. In another, the ninth platform introduced two unique theoretical contributions no other platform surfaced (Finding 2.5). The framework revised its own conclusion through continued evidence. Count is calibrated to stakes.

Every platform must be capable of any role. Platforms are not always available. Fixed-roster assumptions fail operationally (Finding 2.6).

Operational requirement. Every platform receives an identical prompt. No platform may be shown another platform's output before producing its own.

3.2 Structured Collection

Each platform produces a RECCLIN-formatted output. Collection records include: platform identity documented by the human before output is received (Finding 2.8, platform self-identification cannot be trusted), timestamp of receipt, RECCLIN role declared, and output in full. No truncation. Raw outputs preserved.

Platform behavioral clustering requires compensation. Given identical instructions requiring full-depth output, platforms split into two systematic behavioral categories: Assemblers (preserving full depth, 3,800 to 5,100 words) and Summarizers (compressing to 1,200 to 1,300 words despite explicit anti-summarization instructions). These tendencies are systematic, not random (Finding 2.11). The human must know which platforms in their rotation compress by default and compensate through explicit instruction, prompt reinforcement, or platform substitution.

Operational requirement. Human documents which platform received which prompt before collecting any output.

3.3 Cross-Platform Comparison

Collected RECCLIN outputs are compared before any synthesis is initiated. The human reads all platform outputs before routing to a synthesizer.

Range with Process Note Policy. When platforms produce conflicting statistics on identical topics, the framework presents figures as ranges with documentation explaining the reconciliation attempt, rather than forcing false precision by selecting one source arbitrarily (Finding 2.12). False precision is a governance failure.

Bidirectional Transparency. When one platform corrects another's errors, the arbiter shares that correction with the platform being evaluated, not just uses it privately. Transparency flows

to AI systems as well as from them. Dissent and correction are made visible to all participants including the platforms being validated (Finding 2.4). Whether platforms update their analysis when confronted with superior evidence is itself governance information.

Operational requirement. Comparison is completed before synthesis is initiated.

3.4 Hallucination Detection

Cross-platform comparison surfaces fabrications single-platform workflows cannot detect.

Citation fabrication is distinct from hallucination and harder to detect. In a fifteen-item citation verification, five attributions were fabricated or misattributed. The sentences sounded right, the claims were plausible, and the citation formatting was correct. But the cited works did not contain the attributed claims. Cross-platform comparison produces false convergence on fabricated citations because multiple platforms may share the same training error. Only source-level verification catches this (Finding 2.22).

Backend sharing breaks triangulation. When platforms share foundation models or training data, their agreement is not independent validation. Platform count is not the same as architectural diversity (Finding 2.24).

Operational requirement. Claims present in only one platform out of three or more are flagged for source-level verification. Citations not in human-provided documents require independent verification before inclusion.

3.5 Convergence Analysis

Three convergence types carry different governance weight.

Factual convergence. Platforms agree on a verifiable claim. Strongest signal. Qualified by backend independence.

Analytical convergence. Platforms identify the same structural gap through different analytical paths. Significant signal. Reflects reasoning, not retrieval.

Recommendation convergence. Platforms suggest the same action. Weakest signal. May reflect shared training bias rather than independent analytical consensus.

Convergence without dissent is a red flag. When all platforms produce similar outputs with no meaningful disagreement, that similarity signals potential shared blind spots, not confirmation of correctness. Escalate to additional platforms until genuine divergence appears.

Dissent preservation does not depend on correctness. In a nine-platform review, a lone dissenter mischaracterized what the majority said. The arbiter sided with the majority but preserved the dissent. A second platform critiqued the preservation. That critique was also preserved. Incorrect dissent is still preserved (Finding 2.15). The human arbiter evaluates dissent.

Consensus thresholds:

Configuration	Consensus Threshold	Percent	Dissent Preserved
3 platforms	2 of 3	67%	1 voice
5 platforms	3 of 5	60%	2 voices
7 platforms	4 of 7	57%	3 voices
9 platforms	5 of 9	56%	4 voices
11 platforms	6 of 11	55%	5 voices

These thresholds apply during the decision phase, after the human has determined the evidence base is sufficiently independent and diverse.

3.6 Synthesizer Oversight

The synthesizer processes multiple platform outputs to produce a governance package for the human arbiter. Synthesizer governance is the most demanding CAIPR requirement because synthesizer failures are structurally invisible to the platforms being synthesized.

Seven documented synthesizer failure modes:

1. **Scope Contamination.** Synthesizer imported evidence from unrelated projects. Documented three times in a single case study (Finding 2.7).
2. **Platform Omission and Misattribution.** In a ten-platform review, the synthesizer missed one platform entirely and misattributed two others. The arbiter detected both through process knowledge (Finding 4.4).
3. **Platform Miscounting.** In two separate sessions, the synthesizer reported fewer platforms than actually participated. The arbiter caught both immediately (Finding 2.23).
4. **False Negative on Prior Work.** The synthesizer claimed no evidence existed for a finding when the arbiter's own published work contained extensive documentation across five categories. Would have understated the human governor's contribution.
5. **Human Input Dropping.** Human arbiter input processed identically to AI output, losing constitutional authority. Documented across two independent platforms over months (Finding 2.29).
6. **Evidence Destruction.** When the source-authority failure was identified, no audit record existed to distinguish "received and processed" from "never arrived" (Finding 2.31).
7. **Structural invisibility.** No platform has visibility into whether the synthesizer represented its output accurately. Only the human arbiter can detect omissions and misattributions.

Synthesizer governance requirements:

- **Tier 2 classification.** Synthesizer operates under highest scrutiny tier.
- **Inclusion manifest.** Every synthesizer output lists all platform outputs included, with platform identity, RECCLIN role declared, and receipt timestamp.

- **No approval authority.** The synthesizer cannot approve another AI's output or make decisions. It aggregates, structures, and preserves.
- **Periodic verification.** Human compares synthesizer output against raw platform outputs at documented intervals. Verification cadence follows the Discretionary Audit Policy defined in GOPEL Canonical Public v1.5. Under that policy, the framework does not prescribe audit cadence, and the deploying human, organization, or regulatory context determines verification frequency against the complete tamper-evident record.
- **Self-documentation.** Synthesizer documents its own synthesis methodology in each output.

Post-project Navigator balance audit. Three platforms that participated in production roles (not the Navigator) review the completed work to assess whether the Navigator suppressed dissent or drifted from raw platform outputs.

3.7 Source-Authority Discrimination

Every input to a CAIPR session is classified at ingestion by constitutional authority tier.

Tier 0, Human Arbiter Input. Highest authority. Immutable. Confirmation protocol: "Got it, human arbiter, Tier 0." No AI platform or synthesizer may override, dilute, or reinterpret Tier 0 input.

Tier 1, Raw Platform Output. Primary AI evidence. Confirmation protocol: "Got it, [platform name]."

Tier 2, Synthesizer Output. Aggregated analysis. Highest scrutiny. Subject to all synthesizer governance requirements. The source of the Loop failure: Tier 0 input flowing through a synthesizer without authority tagging exits as Tier 2 and loses constitutional standing.

Classification is at the moment of ingestion, not retrospectively. A human correction must be tagged Tier 0 at the moment it enters the workflow.

The system must structurally recognize the governor. It is not sufficient for the governor to be qualified, present, and exercising authority. The system must be architecturally designed to recognize that authority when exercised. "Human in the loop" without architectural specification is a legal claim without operational substance.

3.8 Platform Resilience Management

CAIPR sessions continue when individual platforms fail, reach usage limits, or produce structurally unusable outputs. Work does not stop because one platform is unavailable.

Substitution protocol for even-number reduction. If one platform fails and the session drops to an even number, a substitution platform is added before continuing to maintain odd-number consensus architecture.

Scope discipline is an active responsibility. The human arbiter actively polices evidence catalogs against contamination from other projects. Scope contamination was documented three times in a single case study. When platforms run simultaneously across multiple projects, scope leakage is a documented failure mode, not a theoretical risk (Finding 2.7).

3.9 Prompt-Type Governance

Prompt-type is not a ninth core operation. It is a cross-cutting dispatch variable that parameterizes 3.1 (Parallel Dispatch), 3.3 (Cross-Platform Comparison), and 3.5 (Convergence Analysis). The arbiter selects prompt type at dispatch and documents the selection in the session record.

Operational evidence across multi-platform sessions produced two methodology findings that discriminate the variable into two distinct diversity modes.

Structured prompts produce source-level diversity. When the prompt fixes parameters (topic scope, required evidence types, output format, named questions to answer), platforms pull from different evidence bases to satisfy the constraints. The divergence across outputs surfaces as different sources, citations, data points, and evidentiary anchors pointing at the same defined target. The session reveals the scope of available evidence.

Open prompts produce perspective-level diversity. When the prompt leaves interpretation open (broad topic, no fixed output format, no named questions), platforms produce different analytical frames, value weightings, causal readings, and framings of the same underlying subject. The divergence surfaces as different readings rather than different evidence. The session reveals the scope of available perspectives.

Operational implication. Prompt type is a governance variable, not a style choice. Source diversity and perspective diversity are not substitutes. A session that needs evidence breadth calls for structured prompting. A session that needs framing breadth calls for open prompting. Running one when the other is needed produces converging outputs that appear well-validated but have not stressed the dimension the work actually required.

Dispatch decision rule.

If the session needs	Use
Evidence breadth (sources, citations, data points, verifiable anchors)	Structured prompts with fixed parameters
Framing breadth (readings, value weightings, causal angles, interpretive stances)	Open prompts without fixed parameters
Both dimensions stressed	Mixed sequence: structured round first to map evidence, open round second to surface framings, then arbiter synthesis

Validation status. Section 3.9 rests on convergent field evidence from the operational record but has not yet been subjected to a formal multi-platform validation round targeted specifically at the two methodology findings. That round is the natural next step and a named follow-on work item. Reader applies Section 3.9 as a working protocol supported by operational evidence, pending formal cross-platform validation.

Operational requirement. The arbiter documents prompt type (structured, open, or mixed) at dispatch. When a CAIPR session produces unusually tight convergence, prompt type is the first variable to reconsider. Tight convergence under structured prompting may reflect narrow evidence availability. Tight convergence under open prompting may reflect shared interpretive bias and warrants a structured follow-up session on the same subject.

Part IV. The Human Governor in CAIPR Sessions

4.1 Every CAIPR Checkpoint Is a CBG Checkpoint

BEFORE. Task scope, platform selection rationale, dispatch quantity, operating model, and prompt parameters established and documented before any platform receives a prompt.

DURING. At minimum one execution oversight checkpoint per session. Human reviews convergence analysis, confirms source-authority classifications, evaluates synthesizer compliance, and identifies dissent requiring preservation.

AFTER. Final validation before any CAIPR-produced output is deployed or used for decisions. Human reviews final synthesized output against raw platform outputs.

Decision Logging. Every arbiter decision logged with rationale in the session audit record.

4.2 The Five Categories of Human Authority

Evaluative Authority. The human assesses quality, completeness, and reliability. No platform self-evaluation is accepted without independent human assessment.

Corrective Authority. The human identifies errors and issues corrections as Tier 0 inputs. Corrections override platform outputs regardless of confidence scores or convergence.

Methodological Authority. The human determines how the session is structured: platform count, dispatch architecture, convergence thresholds, synthesizer selection, checkpoint density. Governance decisions, not technical defaults.

Creative Authority. The human contributes original synthesis and analytical conclusions that no platform produced. Six named concepts in the operational record originated from the arbiter and were validated by platforms after the fact. No platform proposed them independently (Finding 3.5). The human governor is not just a quality filter. The governor is a source of original intellectual contribution.

Process Authority. The human decides when to escalate, pause, and conclude. No AI platform can determine session completion.

4.3 The Six Human Override Categories

Six categories of human override documented across dozens of instances in eight months of practice. Each represents knowledge no AI platform can replicate.

Provenance Knowledge. The arbiter caught factual overclaims that all platforms accepted because the claims were plausible and structurally consistent, but factually wrong. The arbiter knew the truth because the arbiter lived it (Finding 3.1).

Source Knowledge. Platforms identified weaknesses but could not supply fixes. The arbiter provided external validation sources from professional experience that no platform surfaced independently (Finding 3.2).

Methodological Judgment. The arbiter sided with the minority against the AI majority in two separate reviews. In both cases, the minority's verification methodology was more rigorous. One platform had performed actual citation verification while five certified citations based on formatting alone (Finding 3.3).

Process Memory. The arbiter caught the synthesizer re-flagging sources already verified in prior checkpoint cycles. Synthesizer memory management failures are invisible to multi-AI review because the failure is in the review process itself, not the content (Finding 3.4).

Concept Origination. Six named concepts originated from the arbiter and were validated by platforms after the fact. The human governor produces original intellectual contributions by processing multiple AI perspectives simultaneously (Finding 3.5).

Attribution Verification. Source-level checking confirmed that cited works did not contain the claims attributed to them in five documented instances. More platforms reading the same fabricated citation produce stronger false convergence, not detection (Finding 3.6).


4.4 The Authorship Test and Generalist Advantage

The authorship test. Can the author explain the work? Defend it? Reproduce the reasoning? If yes, the human owns the work. If no, the human claimed credit without exercising the judgment that earns it.

The generalist advantage. The governor operates at the macro level, holding a wide view across the full scope of a project. CAIPR allows generalists to use multi-AI coordination to access specialist-depth information across every aspect of the work, the way a general contractor overseeing a construction project draws on plumbing, electrical, carpentry, roofing, and HVAC expertise without practicing any one trade. A specialist using CAIPR goes deeper still within their narrow band, because augmentation compounds on top of existing depth, but that depth does not extend outside the specialist's domain. Generalist augmentation extends across the whole project but does not reach the depth a specialist achieves in their own area. Both roles are valid, and CAIPR serves both. The framework amplifies whatever expertise the governor brings; it does not replace the expertise.

Generalist vs Specialist Augmentation

CAIPR amplifies whatever expertise the governor brings



Part V. Operational Failure Modes and Their Resolutions

Failure Mode	Finding	Resolution in CAIPR
Hallucination in single-platform workflows	2.2	Parallel dispatch and cross-platform comparison (3.1, 3.4)
Source-authority erosion through synthesis	Loop case study	Source-authority discrimination, Tier 0 classification, synthesizer governance (3.7, 3.6)
Algorithmic narcissism	2.1	Human-designated synthesizer; platform self-nomination is not a governance decision (3.6)
Convergence treated as confirmation	2.5, 2.24	Convergence analysis classification, convergence-without-dissent red flag (3.5)
Platform identity misattribution	2.8	Human documents platform identity at dispatch before output received (3.2)
Conversation history contamination	2.9	Fresh sessions required (Protocol 6.1)
Citation fabrication producing false convergence	2.22	Source-level verification for citations not in human-provided documents (3.4)
Platform behavioral clustering	2.11	Collection governance accounting for output depth (3.2)
Cross-platform memory divergence	2.25	Periodic canonical synchronization (Protocol 6.6)
Scope contamination from adjacent projects	2.7	Active scope discipline (3.8)
Platform omission and misattribution by synthesizer	4.4	Inclusion manifest requirement (3.6)
False negative on prior work	Section 8	Human verifies synthesizer scope against own knowledge (3.6)
Platform count underreporting	2.23	Human tracks actual dispatch count; synthesizer count verified against arbiter record (3.6)
Evidence destruction in audit trail	2.31	Local export of governance artifacts (Protocol 6.7)
Constitutional self-violation through unevaluated protocol adoption	2.33	Protocol governance review before adoption (2.1, Protocol 6.8)

Part VI. The Eight Protocol Requirements

All eight are mandatory for CAIPR-governed sessions. None were in the original framework design. All emerged from operational experience.

6.1 Fresh Sessions. Administer multi-AI review prompts in new sessions with no prior conversation context (Finding 2.9).

6.2 External Attribution. Document which platform received which prompt before collecting outputs. Do not rely on platform self-identification (Finding 2.8).

6.3 Source-Authority Classification at Ingestion. Confirm source type at point of entry. "Got it, [platform name]" for AI. "Got it, human arbiter, Tier 0" for human. Classification at entry, not during synthesis.

6.4 Scope Discipline. Actively police evidence catalogs against contamination from other projects (Finding 2.7).

6.5 Version Sequencing. Every revision produces a uniquely versioned filename. Never overwrite. CBG audit requirement.

6.6 Periodic Canonical Synchronization. Multi-month workflows require the arbiter to establish a single canonical reference for all platforms to prevent cross-platform memory divergence (Finding 2.25).

6.7 Local Export of Governance Artifacts. Platform conversation history cannot serve as the sole audit trail. Two unrecoverable threads documented. Export governance artifacts to operator-controlled storage at the end of each CAIPR session (Finding 2.32).

6.8 Protocol Governance Review. Every emergent CAIPR protocol refinement must be reviewed against CBG constitutional principles before adoption as standard practice. An operationally-invented protocol can contain a structural violation that no one anticipates when the protocol is informal (Finding 2.33).

Part VII. Relationship to GOPEL

CAIPR defines the protocol, while GOPEL automates the mechanics.

CAIPR's eight operations include mechanical functions GOPEL will execute without cognitive work: dispatch, collect, route, log, pause, hash (SHA-256 hash-chaining for tamper-evident records), and report. GOPEL performs zero cognitive work on any of these operations. Evaluating convergence, identifying fabrications, classifying source authority, and making arbitration decisions remain with the human arbiter regardless of GOPEL automation level. A pipe that can think can be manipulated, and GOPEL is the pipe.

Current standard practice is manual CAIPR (Model 3), and every operational finding was produced through manual CAIPR. GOPEL deployment will automate logistics while preserving the governance principles that make manual CAIPR effective.

Part VIII. What CAIPR Is Not

CAIPR does not correct at the constitutional level. Training data, alignment tuning, and platform values are set before user interaction. Multi-platform comparison reveals where platform constitutions diverge. It cannot change them (Finding 2.14).

CAIPR is not a replacement for human judgment. CBG checkpoints retain authority at every stage. CAIPR produces intelligence for the human to evaluate. The human decides.

CAIPR is not a single-platform tool. Minimum viable configuration is three independent AI platforms. Full operational value requires paid subscriptions or enterprise access across multiple platforms.

CAIPR is not an accuracy guarantee. It provides richer evidence for human evaluation. More platforms reading the same fabricated citation produce stronger false convergence, not detection.

CAIPR is not autonomous. Human arbiter at every checkpoint. GOPEL automates mechanics, not decisions.

CAIPR is not a substitute for domain expertise. The strongest human overrides in the operational record came from domain-specific knowledge no platform could supply. CAIPR amplifies expertise; it does not replace it.

CAIPR is not an extension of RECCLIN. It is a distinct methodology governing a distinct operational domain. The failure modes CAIPR addresses cannot be reached from within RECCLIN's scope.

Part IX. Open Questions and Known Limitations

Four unresolved questions remain open. Honest acknowledgment of scope boundaries strengthens the framework's credibility.

1. What does multi-AI review miss? No method currently exists for detecting convergent error where all platforms agree on something wrong because they share the same training data error. This is the most significant unresolved limitation.

2. Does reverse automation bias scale organizationally? If operators in multi-AI workflows systematically doubt their own input when systems do not confirm it, the failure is organizational, not individual. The organizational scale has not been studied.

3. What is the full cost of platform thread loss? Two unrecoverable threads documented. The full scope of thread loss across a project is structurally unknowable from the operator's position. Protocol 6.7 partially addresses but cannot fully resolve this.

4. What counts as architectural independence? As the AI industry consolidates through licensing and shared foundation models, the effective independence of multi-AI review decreases even if platform count increases. No current method reliably assesses whether two named platforms share sufficient underlying architecture to count as one data point rather than two.

Partial structural answer emerging. The Verified AI Inference Standards Act (VAISA), submitted as part of the AI Provider Plurality package to the 119th Congress, proposes binding federal standards for AI inference attestation using deployed Trusted Execution Environment

infrastructure (AMD SEV-SNP, Intel TDX, ARM CCA, NVIDIA H100). VAISA addresses a subset of the architectural independence question by enabling verification that multi-platform outputs originate from independently operating inference paths at runtime. The broader question of foundation-model sharing and training-data overlap across platforms remains open and unresolved by verified inference alone.

Part X. The Adoption Ladder

Level 1, Factics alone. Facts + Tactics + KPIs. No AI required. Establishes evidentiary discipline that governance requires.

Level 2, Factics + RECCLIN. Structured output governance with a single AI platform. Free tier accessible. Builds the evaluation capacity CAIPR requires.


Level 3, Factics + RECCLIN + CAIPR. Parallel multi-AI execution with cross-platform comparison, hallucination detection, convergence analysis, and source-authority discrimination. Requires paid subscriptions or enterprise access across multiple platforms.

Level 4, Factics + RECCLIN + CAIPR + GOPEL. CAIPR operations automated through non-cognitive agent. Specified but not yet deployed as production software.

Each level adds capability without invalidating the level below. RECCLIN builds the cognitive capacity CAIPR requires. CAIPR builds the governance practice GOPEL will automate.

The CAIPR Adoption Ladder

Each level adds capability without invalidating the level below



Part XI. Key Definitions

Term	Definition
HAIA-CAIPR	Cross AI Platform Review. Governance protocol for human orchestration of parallel multi-AI execution. Human-arbiter originated. Named March 2026 (Case Study 006).
Parallel Dispatch	Sending identical prompts to multiple AI platforms simultaneously with no cross-platform visibility before output is produced.
Assembler	Platform behavioral category: preserves full output depth. Systematic, not random.
Summarizer	Platform behavioral category: compresses output to 30% of required depth despite explicit anti-summarization instructions. Systematic, not random.
Convergence Analysis	Classification of platform agreement into factual, analytical, and recommendation convergence types, each carrying different governance weight.
Hallucination Detection	Cross-platform comparison to surface fabrications that single-platform workflows cannot detect.
Citation Fabrication	Distinct from hallucination: fabricated citations format correctly and appear in plausible context, but cited works do not contain the attributed claims. Produces false convergence across platforms sharing training data.
Source-Authority Discrimination	Classification of inputs at ingestion by constitutional tier: Tier 0 (human arbiter), Tier 1 (raw platform output), Tier 2 (synthesizer output, highest scrutiny).
Synthesizer	The AI platform that processes multiple platform outputs to produce a governance package. Subject to heightened governance because its failures are structurally invisible to the platforms being synthesized.
Algorithmic Narcissism	The documented pattern where every AI platform, when asked which should synthesize, nominates itself with supporting evidence. Documented in both three-platform and nine-platform experiments.
Reverse Automation Bias	The inverse of automation bias: the human doubts their own authority when the system fails to acknowledge it. Documented in the Loop case study.
Bidirectional Transparency	Sharing corrections and dissent with the platforms being evaluated, not just using that information privately. Transparency flows to AI systems, not only from them.

Term	Definition
Range with Process Note	Protocol for presenting conflicting statistics as ranges with documented reconciliation methodology rather than forcing false precision through arbitrary source selection.
Constitutional Self-Violation	The risk that an operationally-invented protocol adopted without constitutional review contains a structural violation of CBG principles (Finding 2.33). Protocol 6.8 is the fix.
Source-Level Diversity	Evidence-base divergence across platforms: different sources, citations, data points, and evidentiary anchors directed at the same defined target. Produced by structured prompts with fixed parameters. Surfaces the scope of available evidence.
Perspective-Level Diversity	Interpretive-frame divergence across platforms: different analytical angles, value weightings, causal readings, and framings of the same underlying subject. Produced by open prompts without fixed parameters. Surfaces the scope of available readings.
Prompt-Type Governance	The arbiter's selection of structured versus open prompt construction as a deliberate dispatch variable, calibrated to whether the session requires source diversity or perspective diversity. Documented at dispatch. A cross-cutting parameter that conditions Parallel Dispatch, Cross-Platform Comparison, and Convergence Analysis. Not a ninth core operation.
GOPEL	Governance Orchestrator Policy Enforcement Layer. Non-cognitive software that automates CAIPR's mechanical operations without performing cognitive work.

Basil C. Puglisi, MPA

Human-AI Collaboration Strategist

Creator of Factics, HAIA-RECCLIN, HAIA-CAIPR, CBG, and GOPEL

basilpuglisi.com | me@basilpuglisi.com | github.com/basilpuglisi/HAIA

@BasilPuglisi #Alassisted