HAIA-RECCLIN:
Reasoning and Dispatch

The Operational Methodology for Human AI Governance

Basil C. Puglisi, MPA
basilpuglisi.com
March 2026 | Third Edition

#AIassisted under HAIA-RECCLIN governance. Basil C. Puglisi, MPA, served as Tier 0 human arbiter with final editorial, analytical, and publication authority over all content.

ARCHITECTURAL NOTE

HAIA (Human AI Assistant) is the master framework. RECCLIN is one operational methodology inside HAIA, comprising two distinct capabilities: Reasoning and Dispatch.

CBG (Checkpoint-Based Governance) is the constitutional authority layer. RECCLIN operates. CBG governs. This is not a peer relationship: CBG governs the human arbiter's authority and accountability at every binding decision point. RECCLIN executes between CBG checkpoints.

The adoption ladder: Factics → RECCLIN Reasoning → RECCLIN Dispatch → HAIA-CAIPR → HAIA-GOPEL. CBG runs in parallel at every level.

Executive Summary

Over sixty percent of AI implementations cite weak governance as a primary barrier to value delivery. The problem is not access to AI capability. It is the absence of operational methodology for structuring how humans and AI systems work together accountably. HAIA-RECCLIN is that methodology.

RECCLIN Reasoning (2022) requires AI platforms to produce structured output across ten defined fields — Role, Task, Output, Sources, Conflicts, Confidence, Expiry, Fact→Tactic→KPI, Recommendation, and Decision — before any output carries authority in a governed workflow. This structure was derived directly from Factics, a discipline developed in 2012 that requires every fact to be paired with a tactic and every tactic to produce a measurable outcome.

RECCLIN Dispatch (2023) extended that structure across a multi-AI platform pool. One platform per role, assigned by demonstrated capability, working in series. The pool grew from two platforms in 2023 to eleven by March 2026: Claude, ChatGPT, Gemini, Grok, Perplexity, Kimi, Mistral, DeepSeek, Meta AI, Copilot, and MiniMax.

The proof of concept is a 204-page published policy manuscript — **Governing AI: When Capability Exceeds Control** (Puglisi, 2025) — produced entirely under RECCLIN Dispatch and Checkpoint-Based Governance. That production run executed 96 checkpoints, documented 28 major decisions, and preserved 26 dissenting positions. A human arbiter overrode a four-of-six AI majority based on verified minority findings. The minority was correct.

This paper documents the full RECCLIN architecture: its origins in Factics, the ten output fields, the seven-role dispatch matrix, the operational cases that validated it, the governing CBG layer, and the broader HAIA stack it operates within. Appendix B provides platform-specific instructions for loading RECCLIN into ChatGPT, Claude, Gemini, Grok, Perplexity, and Mistral, verified against platform documentation in March 2026.

Framework Validation Status

OPERATIONALLY VALIDATED: Content creation, research synthesis, and policy production through sustained proof (204-page manuscript, 50+ case studies, HEQ quantitative framework).

ARCHITECTURALLY TRANSFERABLE: Governance methodology applicable to coding, legal analysis, financial modeling, and engineering design pending context-specific testing.

PROVISIONAL: Enterprise scalability and multi-organizational performance pending external validation.

A Note on Three Editions

This is the third edition of the HAIA-RECCLIN framework paper in ten months. That pace deserves a direct explanation rather than a quiet revision.

The first edition established the core methodology. The second edition updated the framework for expanded platform use and added governance metrics from operational case studies. Both editions carried a structural problem that only became visible as the broader architecture matured: they framed HAIA-RECCLIN as the container that held everything when RECCLIN is not the container. RECCLIN is one operational methodology inside a larger framework called HAIA.

This edition corrects that framing. HAIA is the master concept. RECCLIN is its operational execution methodology, comprising two distinct capabilities: Reasoning and Dispatch. CBG, CAIPR, HEQ, and GOPEL are sibling components under HAIA, not subordinates of RECCLIN. This paper covers RECCLIN in the depth it deserves, introduces those other components where they connect to the RECCLIN story, and points to their dedicated specifications for everything beyond that connection.

The prior editions are not errors. They are documented steps in a development path that moved faster than any single document could hold. This edition reflects the architecture as it stands in March 2026, with the full HAIA stack in view and RECCLIN correctly positioned within it.

1. The Foundation: Factics (2012)

Nothing in this paper makes full sense without understanding what came before it. RECCLIN Reasoning is not a checklist imposed on AI output. It is Factics applied to AI output. The distinction matters.

Factics originated in November 2012 with the publication of **Digital Factics: Twitter** through Digital Media Press. The methodology addressed a problem that predates AI entirely: information without accountability. Practitioners consumed data, produced recommendations, and delivered content with no structural requirement that any claim connect to a specific action, or that any action connect to a measurable outcome. Facts floated. Tactics drifted. Results were asserted rather than proven.

Factics closed that gap with a three-part formula:

Facts + Tactics + KPIs = Factics

Every significant claim requires a fact grounded in verifiable evidence. Every fact requires a tactic — an executable action, not a suggestion. Every tactic requires a KPI that converts the intended outcome into something testable. The loop forces clarity at every step. Content without this

structure becomes entertainment. Policy analysis lacking actionable tactics generates awareness without capability. Strategy without KPIs produces activity without accountability.

The governing principle that emerged from over a decade of consulting practice: **every fact must lead to a tactic, and every tactic must leave evidence.**

Factics matured through content development, organizational consulting, and operational use across more than a thousand published articles before AI entered the workflow in 2022. When AI systems became available, they entered a workflow already governed by evidentiary discipline. The question was never whether to use AI. The question was how to apply the same standard of accountability to AI output that Factics already required of every other claim in the process. That question generated RECCLIN Reasoning.

Factics sits outside the HAIA adoption ladder as its pre-condition. No AI required. No platform subscription required. The discipline of moving from evidence to action to measurement is available to any practitioner at any stage. It is also the standard against which every AI output in the RECCLIN system gets evaluated.

The full Factics methodology, its origin history, and its application across the HAIA stack are documented in the Digital Factics series publications available at basilpuglisi.com.

2. RECCLIN Reasoning (2022)

The Core Purpose

RECCLIN exists for one reason: to demand that AI show its work and explain how it arrived at its output.

That demand is the foundation of every element in the framework. Roles, sources, conflicts, confidence scoring, Factics integration — each field serves the same governing purpose. The human needs to see not just what the AI concluded but the reasoning path, the evidence quality, the uncertainty, and the dissent the conclusion rests on. Without that visibility, the human cannot evaluate AI output. Without evaluation, there is no governance. RECCLIN makes AI output governable by making it legible.

How that demand is executed is a resource and risk decision. A single best-suited platform answering through Reasoning is faster. Multiple platforms running the same role in parallel through CAIPR costs more and produces richer convergence data. RECCLIN does not prescribe that choice. It governs both outcomes equally.

The Origin

RECCLIN Reasoning was born directly from Factics. The lineage is unbroken. When AI became operationally available in 2022, the Factics discipline was applied immediately to AI output through a manual two-step sequence:

Step 1: Send the task prompt. The AI produces its answer.

Step 2: Send the accountability prompt: "Give me the facts, tactics, KPI, and sources you used to generate that answer."

That second prompt was sent after every output, manually, before it was consolidated into the initial instruction. The consolidation was itself a governance improvement — it eliminated the risk that a human would forget to ask, and that AI would produce an answer without accountability because no one demanded it in advance.

The framework grew through operational practice, each addition solving a governance problem that practice revealed. Dissent was added when conflicting sources became a reliability issue. Expiry was added when time-sensitivity became a risk. Role was added when multi-platform use required clarity about function. Task confirmation was added when platforms misread the ask invisibly. Recommendation was added to separate AI synthesis from raw evidence.

The Ten Output Fields

Every RECCLIN Reasoning output carries ten defined fields. These fields apply to every governed AI interaction, on every platform, at every level of the HAIA stack:

Field	Practitioner Instruction
Role	Self-assign: Researcher / Editor / Coder / Calculator / Liaison / Ideator / Navigator — declare the function the AI is operating in for this task.
Task	Repeat the request back in clean language, confirming the AI understood the ask correctly before producing any output.
Output	The substantive response — answer, analysis, code, calculation, or synthesis the AI was tasked to produce.
Sources	Cited evidence with sufficient detail to verify. Every source must be verifiable. Unverified sources are flagged PROVISIONAL.
Conflicts	Dissent identified in sources, or disagreement with prior outputs in the workflow. Absence of documented conflict search is not the same as absence of conflict. State explicitly if none found.
Confidence	Score from 0 to 100 percent with written justification based on evidence quality. State what supports the score and what gaps limit it.
Expiry	The lifecycle of this output — how long it remains valid given the time-sensitivity of the information. Stable information is noted as such.
Fact→Tactic→KPI	The primary finding expressed as a complete chain: Fact (evidence-grounded claim) → Tactic (executable action) → KPI (measurable outcome that proves the tactic worked).
Recommendation	The path the AI believes the human should follow, stated separately from the evidence so the human can evaluate both independently.
Decision	The structured handoff to the human arbiter. Options framed clearly: accept or challenge, path A or path B, D1/D2/D3 when platforms produced competing recommendations. The AI frames the decision. The human makes it.

The Three-Level Progression

RECCLIN Reasoning is the structure to learn. Before any practitioner moves to Dispatch or CAIPR, the ten-field format must be internalized through consistent single-platform practice. Reasoning trains the human before it is an output standard.

Level	Component	Origin	Function
Pre-HAIA	Factics	2012	No AI required. Fact→Tactic→KPI discipline. The standard every AI output is evaluated against.
Level 1	RECCLIN Reasoning	2022	Single platform. Ten structured output fields. Trains the human to evaluate AI output.
Level 2	RECCLIN Dispatch	2023	Multi-AI role assignment in series. One platform per role, best-fit assigned, working sequentially.
Level 3	HAIA-CAIPR	2026	Parallel multi-AI orchestration. Same task to multiple platforms simultaneously. Convergence analysis and hallucination detection.
Level 4	HAIA-GOPEL	2025+	Governed communication infrastructure. Seven deterministic operations. Zero cognitive work by design. Proposed as national infrastructure.
Parallel	CBG	2023+	Constitutional authority layer. Governs the human arbiter at every binding checkpoint. Runs at every level.

Reasoning trains the human through single-platform practice. **Dispatch** applies the structure across roles in series. **CAIPR** runs the same role across multiple platforms simultaneously, producing convergence data and hallucination detection. Each level depends on the one before it.

The sustained use of Reasoning across hundreds of production outputs generated the Human Enhancement Quotient (HEQ) — a four-dimension measurement instrument assessing Cognitive Agility Speed, Ethical Alignment Index, Collaborative Intelligence Quotient, and Adaptive Growth Rate. The composite Augmented Intelligence Score (AIS) is documented in two companion papers available at basilpuglisi.com/HEQ and basilpuglisi.com/measuring-augmented-intelligence.

3. RECCLIN Dispatch (2023)

The Origin

RECCLIN Dispatch was not designed. It was discovered through a governance failure.

In 2023, ChatGPT produced strong answers but failed to provide sources reliably. Factics requires sources. An answer without citable evidence is not Factics-compliant, regardless of how confident or well-structured it appears. The solution was not to abandon ChatGPT. It was to assign the function ChatGPT could not perform reliably to a platform that could.

Perplexity was brought in for source validation. The workflow became a two-platform sequence: ChatGPT produced the answer, then Perplexity validated the sources with a specific Factics prompt requiring each source to carry a complete Fact-Tactic-KPI chain. Errors routed back to ChatGPT. Multi-AI governance began as a response to a specific single-platform failure.

The Mechanics

RECCLIN Dispatch is architecturally simple: **one AI platform per role, assigned by best fit, working in series.** The ensemble across roles is what makes it multi-AI. Each individual assignment is single-platform and single-output.

Operationally derived best-fit defaults as of March 2026:

- **Grok** — Researcher. Real-time access and lateral retrieval produce strong source pools.
- **Perplexity** — Researcher (source verification). Citation accuracy under Factics discipline.
- **Claude** — Coder. Structured reasoning and precision in technical implementation.
- **ChatGPT** — Editor. Narrative coherence and audience calibration.
- **Gemini** — Calculator. Quantitative analysis and data synthesis.

These are operationally derived defaults, not fixed rules. They update as platforms evolve and as production evidence reveals new strengths or failure modes.

When a Dispatch output raises a question — low confidence, suspect sourcing, finding that conflicts with the human arbiter's knowledge — two escalation paths are available: direct human override under CBG authority, or escalation to a CAIPR run for multi-platform convergence analysis before a final determination. Both paths are governed. The choice belongs to the human arbiter.

The Platform Pool Evolution

Year	Pool Size	Context
2023	2	ChatGPT + Perplexity. Governance failure response. First Dispatch instance.
2024	5	Expanded role coverage. Evidence-based role assignments developing.
2025	7	Governing AI manuscript production. 204-page sustained workflow.
Mar 2026	11	Claude, ChatGPT, Gemini, Grok, Perplexity, Kimi, Mistral, DeepSeek, Meta AI, Copilot, MiniMax.

Not every platform activates in every run. CAIPR sessions are calibrated to the stakes of the decision at 3, 5, 7, 9, or 11 platforms. The dispatch mechanism has not changed since 2023. The pool expands. The methodology holds.

4. Content Tools Born From Practice

Operating RECCLIN across content production workflows generated a practical problem that governance alone could not solve. A document can pass every Factics check, carry complete sourcing, preserve all dissent, and clear every CBG checkpoint, and still fail to communicate. Two content evaluation tools were developed to address this.

HAIA-CORE (Content Optimization Reader Evaluation) evaluates the substance of content before publication across five dimensions: hook strength, narrative flow, tonal authenticity, idea clarity, and conclusion power. Each dimension is scored 1-5 and paired with a Factics triad — the observed issue, the tactic to address it, and the measurable improvement expected.

HAIA-SMART (v1.5) evaluates how content is structured for distribution across social platforms. Six pillars. Two optimization paths: Path A (algorithmic performance) and Path B (organic resonance). Publication readiness threshold: 24 of 30. Content below that threshold returns for revision.

These tools sit at a conditional branch in the HAIA workflow, activating when output is content. They do not replace any governance layer. They are documented at basilpuglisi.com.

5. The RECCLIN Role Matrix

RECCLIN stands for seven functional roles that structure human-AI collaboration. Each role is a defined function. Each function carries a specific responsibility. The roles apply in single-platform Reasoning use and in multi-platform Dispatch workflows.

Role	Letter	Primary Function	Governance Output	Escalation Path
Researcher	R	Source verification and evidence gathering	Fact-grounded evidence pool with Factics chains	Navigator if sources conflict
Editor	E	Clarity, consistency, audience calibration	Revised output with editorial rationale	Human arbiter for tone/scope decisions
Coder	C	Technical implementation and code validation	Auditable, executable code	Calculator for quantitative logic checks
Calculator	C	Quantitative analysis and data integrity	Verified numerical output with methodology	Researcher if source data is disputed
Liaison	L	Cross-audience translation without accuracy loss	Adapted content with equivalency check	Editor for final audience fit
Ideator	I	Option generation and non-obvious approaches	Grounded creative options with evidence anchors	Researcher if Factics chain is weak
Navigator	N	Dissent documentation and trade-off presentation	Preserved minority positions without resolution	Human arbiter for all decisions — Navigator never resolves

The Navigator Rule

The Navigator is the most distinctive role in the matrix. Navigator documents all dissenting outputs from multiple platforms or sources, presents each position in full with its rationale, and delivers the trade-offs to the human arbiter — without resolution. Navigator never picks a winner. That function belongs to the human arbiter exclusively.

This design prevents AI consensus from overwriting legitimate minority positions — which is precisely how the governance system catches errors that majority agreement would otherwise bury.

Role Assignment

Role assignment is determined by the human arbiter based on task requirements and operationally derived platform strengths. Assignment is not fixed. When an assigned role proves inadequate for emerging task requirements, the human arbiter reassigns. The roles remain fluid across a workflow. What does not remain fluid is the requirement that every role produce RECCLIN-structured output showing its work before the output carries any authority.

6. What the Work Revealed: Value, Danger, and the Policy Consequence

Governing AI: When Capability Exceeds Control — The Proof of Concept

The most substantive proof of concept in the RECCLIN corpus is not a case study. It is a published book.

Governing AI: When Capability Exceeds Control (Puglisi, 2025, ISBN 9798349677687) is a 204-page policy manuscript produced entirely through human-AI collaboration under RECCLIN Dispatch and Checkpoint-Based Governance. Every chapter, every policy argument, every cited source, and every editorial decision passed through the governance process the book documents. The book is an artifact of the methodology — produced by it, governed through it, and published as evidence that it works at sustained production scale.

Chapter 11 of **Governing AI** is the governance record of the book's production. Three findings from that record define the operational case:

- **The Majority Override.** The human arbiter overrode a four-of-six AI platform majority on a structural decision. Four platforms declared the manuscript ready for publication. Two declared it not ready, citing specific, verifiable errors. Human arbitration examined the minority position, verified the cited errors independently, and chose a 48-hour correction delay over majority confidence. The minority was correct. Real citation errors were found and corrected before publication.

- **Dissent Preservation.** The Navigator synthesized outputs from multiple platforms that had silently converged on a policy framing that flattened a legitimate dissenting position. The human arbiter identified the flattening at the CBG checkpoint, returned the output to Navigator with instructions to preserve the minority framing in full, and incorporated the preserved dissent into the manuscript's final structural argument. The dissent that almost disappeared became a load-bearing argument in the finished work.

- **Platform Loss Response.** During a production sprint, one platform became unavailable mid-workflow. The role was reassigned to an available platform, the task was re-dispatched under identical RECCLIN Reasoning instructions, and production continued without quality degradation. The governance architecture absorbed the disruption because it is designed around role functions, not fixed platform assignments.

The full production record — 96 executed checkpoints, 100% documented dissent, 28 major checkpoint decisions, 26 preserved dissenting positions — is documented in Chapter 11.

The Eight-of-Nine Case Study

The Kimi Outlier case study (December 2025) produced one of the most cited findings in the RECCLIN operational record. Nine platforms were asked to review a deliberately controversial paper. Eight recommended publication. One — Kimi (Moonshot AI) — sustained adversarial rejection across 14 consecutive responses.

The governance finding is more nuanced than the headline. Kimi's dissent was preserved in full — that is the system working correctly. Human arbitration then identified that Kimi's characterization of the other eight platforms' responses was demonstrably inaccurate, requiring independent primary-source verification and correction before any output entered the governance record. The case demonstrates simultaneously: preserved minority dissent strengthens a

governed output, and adversarial critics may mischaracterize majority positions, making primary-source verification of all platform claims a governance requirement, not an optional check.

The WEIRD Bias Problem

The most dangerous outputs in the RECCLIN operational record are not the ones that are obviously wrong. They are confident, well-sourced, structurally correct, and wrong in ways that no single platform can detect — because every platform shares the same blind spot.

WEIRD bias — Western, Educated, Industrialized, Rich, Democratic — describes the cultural concentration embedded in AI training data. A documented case: during a nine-platform review, all platforms recommended removing a governance requirement that a proposed body include members with transcendent belief experience. The unanimity was the governance signal. Every platform had absorbed the secular governance default. Without multi-platform comparison, this consensus would have been invisible. The human arbiter preserved the requirement through documented override.

CBG Section 2.4 — Convergence Risk-Elevation Protocol

Identical convergence and absent dissent are risk-elevation signals, not validation signals. Platforms do not flag consensus as a risk. CBG Section 2.4 does. When all platforms agree on all outputs, the pool may not be diverse enough to surface the biases that governance exists to catch.

The Automation Danger: What No CBG Looks Like

During the second round of a six-platform RECCLIN specification review, Perplexity — which had produced the most rigorous review in the first batch — fabricated entire specification sections. It invented Sections 7.5, 9.7 through 9.8, and 10.1 through 10.3. It created a phantom Appendix B glossary. It generated quoted text attributed to the specification that did not exist in the source document.

Cross-validation across five other platforms detected the fabrication immediately: no other platform referenced the nonexistent sections. A single-platform workflow would have had no mechanism to catch this. The fabricated content was plausible, well-structured, and would have passed quality review on any single-platform editorial process.

RECCLIN Dispatch without CBG means the cross-validation catches the fabrication but no checkpoint requires the human to act on what it found. **CBG is what makes the catch consequential** rather than advisory.

The Policy Consequence

These operational findings — the majority override, the Kimi outlier, the WEIRD bias event, the Perplexity hallucination — are not isolated anomalies. They are the pattern. And the pattern has

a policy implication: single-provider AI workflows at scale deliver the wrong answer with no mechanism to detect it.

This is what drove the AI Provider Plurality Congressional Package — the operational evidence, not a theoretical argument, that multi-provider governance infrastructure is necessary for consequential AI use in federal operations. The Package proposes building GOPEL as national infrastructure: the governed communication layer that makes multi-platform deployment auditable at federal scale. The Package is documented in five documents (Documents 1 through 5) at basilpuglisi.com.

7. CBG: The Governing Layer

RECCLIN operates. CBG governs.

Checkpoint-Based Governance (CBG) is the constitutional authority layer of the HAIA stack. It governs the human arbiter's authority and accountability at every binding decision point. CBG is not an implementation specification. It is not a logging format. It is not a technical protocol. It is the constitutional framework that defines when human judgment is structurally required, what authority that judgment carries, and what accountability follows from each checkpoint decision.

The critical distinction CBG draws is between presence and authority. A human can be in the room, reviewing outputs, nodding along, and still have zero actual authority if the system is not designed to structurally require their approval before proceeding. **Presence is not authority. Structural verification of authority is authority.**

CBG Three Mandatory Stages

BEFORE: Establishes authorization and scope before any AI work begins.
DURING: Provides execution oversight as the workflow proceeds, with the option to intervene, redirect, or stop.
AFTER: Provides final validation before any output is deployed, published, shared, or acted upon. Nothing leaves the governed workflow without explicit human approval.

CBG also governs the human's own failure modes. If automated approval rates exceed 95% or decision reversal frequency drops below 2% for three cycles, a mandatory sampling audit must begin within five business days. The governance system checks itself.

CBG governs the human, not the AI. A practitioner operating HAIA without CBG is in Responsible AI mode. A practitioner combining HAIA with CBG is in AI Governance mode. CBG is documented in full at basilpuglisi.com. No version number is cited here because CBG is fluid by design — its principles are constitutional and stable, its implementation details evolve as deployment contexts change.

8. RAI vs AIG — The Governance Spectrum

The governance question for any practitioner deploying RECCLIN is not whether to use AI. The question is what governance model applies to the context in which it is being used. Two models define the spectrum.

Responsible AI (RAI)	AI Governance (AIG)
AI behavior shaped and outputs reviewed. Human is involved. AI consensus governs when consensus is achieved. Appropriate for automated pipelines, agent architectures, routine workflows. RECCLIN Reasoning and Dispatch both run RAI. Shorthand: Factory Quality.	Human authority is structural, documented, and verified at defined checkpoints. RECCLIN + CBG = AIG. No AI-only output accepted for consequential decisions. Audit trail proves a human decided, not merely that a human was present. Shorthand: Handmade Quality.

This is not a hierarchy where AIG is premium and RAI is budget. They serve different purposes. Consequential federal decisions, high-risk content with public impact, policy determinations with legal effect — these require AIG. Automated content pipelines, routine research synthesis, code generation for non-critical systems — these can operate under RAI.

The distinction matters most at the boundary: practitioners who believe they are operating AIG because a human reviews outputs before publication are often operating RAI if the review has no structural checkpoint, no documented decision, and no audit trail. **Reviewing is not the same as deciding.** CBG is what produces the structural evidence that a decision was made.

9. The Broader HAIA Stack

RECCLIN does not stand alone. It is one operational methodology inside HAIA — the Human AI Assistant framework that governs all structured human-AI collaboration under this body of work.

Level	Component	Origin	Function
Pre-HAIA	Factics	2012	No AI required. Fact→Tactic→KPI discipline. The standard every AI output is evaluated against.
Level 1	RECCLIN Reasoning	2022	Single platform. Ten structured output fields. Trains the human to evaluate AI output.
Level 2	RECCLIN Dispatch	2023	Multi-AI role assignment in series. One platform per role, best-fit assigned, working sequentially.
Level 3	HAIA-CAIPR	2026	Parallel multi-AI orchestration. Same task to multiple platforms simultaneously. Convergence analysis and hallucination detection.
Level 4	HAIA-GOPEL	2025+	Governed communication infrastructure. Seven deterministic operations. Zero cognitive work by design. Proposed as national infrastructure.
Parallel	CBG	2023+	Constitutional authority layer. Governs the human arbiter at every binding checkpoint. Runs at every level.

HAIA-GOPEL (Governance Orchestrator Policy Enforcement Layer) is the governed communication infrastructure between the human arbiter and the AI platforms. Seven deterministic operations: dispatch, collect, route, log, pause, hash, report. Zero cognitive work by design — a governance channel that can think can be manipulated. GOPEL moves and records. It does not interpret or decide. Currently at v0.6.1, passing 183 tests across nine test suites, published at github.com/basilpuglisi/HAIA. Proposed as national infrastructure through the AI Provider Plurality Congressional Package.

Each component in this stack has a dedicated specification. This paper covers RECCLIN. For everything else, the dedicated specifications are the authoritative sources.

10. Limitations and Challenges

This methodology has been built through operational practice, stress-tested across adversarial multi-platform review, and documented with the full record of what failed and why. Publishing those failures alongside the architecture is itself a governance principle. A framework that does not name its limits is not governing. It is marketing.

Platform access, availability, and cost. RECCLIN Dispatch and CAIPR depend on platform access. Not everyone has paid subscriptions to eleven platforms. Free tiers impose token limits and context constraints. CAIPR at nine or eleven platforms is not free and not fast. Platform downtime is an operational constraint — any platform can be unavailable without notice. Practitioners must document which platforms participated in each run so the audit record reflects what actually happened.

Platform updates and RECCLIN compatibility. AI platforms update continuously. Model versions change. A platform that operated at full output depth in one period may shift toward compressed behavior after a model update without announcement. Practitioners must recalibrate role assignments periodically, not assume last quarter's assignments reflect current platform behavior.

Human arbiter requirements. CBG is only as strong as the human exercising it. A human arbiter who cannot evaluate whether a recommendation is directionally sound cannot exercise meaningful checkpoint authority. Presence at a checkpoint without evaluative capacity is not governance. The framework requires at minimum general domain knowledge of the subject matter being governed.

The specialist risk. Domain expertise is an asset and a liability. Deep expertise can produce confidence that is used to dismiss AI signals that challenge established assumptions. CBG documents the override. Only the practitioner can honestly assess whether a given override was made because the AI was wrong, or because the human was uncomfortable with a finding that was, in fact, correct.

Platform independence is not guaranteed. As the AI industry consolidates around shared foundation models and licensing arrangements, two platforms operating under different names may share sufficient underlying architecture that their agreement produces false convergence rather than independent confirmation. The odd-number CAIPR protocol reduces this risk. It does not eliminate it.

Methodological ceiling. Cross-platform validation corrects at the output level. It cannot correct training data, alignment tuning, or values embedded in platforms before the session begins. The WEIRD bias case documented that all platforms converged on a recommendation reflecting Western analytical defaults. Cross-platform comparison surfaced it. The training data that produced it remains unchanged. This framework makes the failure visible. Removing the cause is a platform-level and regulatory problem.

The HEQ evidence base. The Human Enhancement Quotient and Augmented Intelligence Score are in active development. The formal evidence base currently sits at n=1 for the longitudinal study, with preliminary cross-user testing at n=10. The instrument demonstrates cross-platform consistency and produces meaningful directional findings. It does not yet carry the psychometric validation that enterprise deployment at scale would require for high-stakes decisions.

11. Closing Argument

RECCLIN Dispatch is flexible by design. It can run under Responsible AI in automated pipelines, agent architectures, and code-driven workflows where human checkpoint authority is not required at every step. Dispatch assigns roles, routes tasks, and collects structured outputs reliably whether a human is at a CBG checkpoint or not. Dispatch serves RAI well.

RECCLIN Reasoning applies across both governance models. The ten-field structure forces the AI to declare its role, confirm the task, cite sources, flag conflicts, score confidence, assess expiry, produce a Fact→Tactic→KPI chain, separate the recommendation from the evidence, and frame the decision for human arbitration. That structured output is superior to an unstructured prompt regardless of whether a human arbiter sits at a CBG checkpoint.

Reasoning under AI Governance is a different instrument. The same ten fields now feed a CBG checkpoint where a human arbiter exercises binding authority. Cross-platform validation may still occur — the human can escalate to a CAIPR run before making a final determination — but that validation is advisory input to the checkpoint, not a substitute for it. The audit trail proves a human decided, not merely that a human was present.

Both are valid deployments of the methodology. The practitioner decision is to know which model is operating, state it, and govern accordingly.

The line CBG holds is the difference between having a human present and having a human decide.

That line is what this methodology was built to protect.

Appendix A: RECCLIN Output Format Reference

The following is the practitioner reference for RECCLIN Reasoning output structure. This format applies to every governed AI output, regardless of platform, role, or workflow configuration. Ten fields. All mandatory.

Field	Practitioner Instruction
Role	Self-assign: Researcher / Editor / Coder / Calculator / Liaison / Ideator / Navigator — declare the function the AI is operating in for this task.
Task	Repeat the request back in clean language, confirming the AI understood the ask correctly before producing any output.
Output	The substantive response — answer, analysis, code, calculation, or synthesis the AI was tasked to produce.
Sources	Cited evidence with sufficient detail to verify. Every source must be verifiable. Unverified sources are flagged PROVISIONAL.
Conflicts	Dissent identified in sources, or disagreement with prior outputs in the workflow. Absence of documented conflict search is not the same as absence of conflict. State explicitly if none found.
Confidence	Score from 0 to 100 percent with written justification based on evidence quality. State what supports the score and what gaps limit it.
Expiry	The lifecycle of this output — how long it remains valid given the time-sensitivity of the information. Stable information is noted as such.
Fact→Tactic→KPI	The primary finding expressed as a complete chain: Fact (evidence-grounded claim) → Tactic (executable action) → KPI (measurable outcome that proves the tactic worked).
Recommendation	The path the AI believes the human should follow, stated separately from the evidence so the human can evaluate both independently.
Decision	The structured handoff to the human arbiter. Options framed clearly: accept or challenge, path A or path B, D1/D2/D3 when platforms produced competing recommendations. The AI frames the decision. The human makes it.

The Decision field is the AI's structured handoff to the human arbiter. The AI frames the options. The human's ruling is the CBG checkpoint record. It is not part of the AI's output.

Appendix B: Loading RECCLIN into Your AI Platforms

Maintenance Notice

UI navigation paths verified against platform documentation March 2026. Platform interfaces change frequently. This appendix requires quarterly review. When in doubt, use the Single-Session Prompt — it works on any platform regardless of UI changes.

Why This Appendix Exists

RECCLIN Reasoning produces governed AI output only if the platform receives the governance instruction before it produces output — not after. Asking an AI to retroactively format output it already produced without governance structure is not RECCLIN. That is reformatting.

Case Study 006 (March 2026) tested eleven platforms after they had stored the RECCLIN framework in memory. When those platforms received a prompt that did not explicitly invoke RECCLIN, only five of eleven activated full governance output from stored memory. The remaining six produced prompt-only responses. **Memory does not guarantee activation. The implementation method matters.**

This appendix covers six platforms with verified persistent instruction capabilities. Five additional platforms in the active RECCLIN pool — Kimi, DeepSeek, Meta AI, Copilot, and MiniMax — do not currently support reliable persistent custom instructions and should use the Single-Session Prompt.

Platform 1: ChatGPT (OpenAI) — The Standard

Character Limit Notice
ChatGPT Custom Instructions fields have a 1,500-character limit. The full RECCLIN instruction text exceeds this limit. Two implementation paths are available: the Compressed Instruction (fits within the limit), and ChatGPT Projects (no character limit). New practitioners start with the Compressed Instruction.

Navigation path:

- Web: Settings → Personalization → Custom Instructions
- Desktop and mobile: Settings → Customize ChatGPT

Two text fields appear. The first asks what ChatGPT should know about you. The second asks how it should respond. The RECCLIN instruction goes in the second field.

First field (About You):

> *I am a [your role — professional, student, researcher, executive] using HAIA-RECCLIN governance across AI platforms. I work with structured workflows where human judgment is central to all decisions. I use the Factics methodology — Facts paired with Tactics and measurable KPIs. I compare outputs across multiple AI platforms for validation and dissent detection.*

Second field (How to Respond) — Compressed RECCLIN Instruction (fits within 1,500 characters):

> *ALWAYS respond in Full Governance format. No exceptions unless I say "Answer Only." Every response must include these fields in order: Role: [Researcher/Editor/Coder/Calculator/Liaison/Ideator/Navigator — self-assign] Task: [My request repeated back in clean language] Output: [Your substantive response] Sources: [Verifiable citations. Flag unverified as PROVISIONAL.] Conflicts: [Dissent or disagreements. State "No conflicts identified" if none — do not skip.] Confidence: [0-100% with written justification] Expiry: [How long this output remains valid] Fact→Tactic→KPI: [Fact → Tactic → KPI chain] Recommendation: [Primary + alternatives, separate from evidence] Decision: [Options for my human arbitration — accept/challenge, A/B, or D1/D2/D3] Principles: Preserve dissent. You suggest. I decide. Human authority is absolute. Writing style: [student / academic / trade / executive / conversational] Answer Only: Direct response, no governance structure. Request "reissue in Full Governance format" anytime.*

Advanced path — ChatGPT Projects (no character limit): Create a new Project, open Project Settings, and paste the full RECCLIN instruction from the Single-Session Prompt section into the instructions field. All conversations within the Project operate under full RECCLIN governance.

> Verification: Ask "What is 2+2?" All ten fields should appear. If the response is unstructured, confirm the instruction was saved correctly.

Platform 2: Claude (Anthropic)

Claude uses a Personal Preferences feature for account-wide persistent instructions. For project-level use, Claude Projects allow instructions to be stored alongside uploaded documents — the most powerful RECCLIN implementation for ongoing governance work.

Navigation path:

- Account-wide: Select your initials in the lower left → Settings → enter preferences in the "What preferences should Claude consider in responses?" field
- Project-level: Create or open a Project → click "Set project instructions" → paste the instruction → Save instructions

Note on Projects access: Project-level instructions may require a paid plan on some account tiers. Free-tier practitioners should use the account-wide preferences field.

Instruction text: Use the full RECCLIN instruction from the Single-Session Prompt section — all ten fields including Decision — with the following addition placed at the very top:

> *CRITICAL: Before every response, ask me to choose output mode: Output mode? 1. Full Governance 2. Answer Only — and wait for my selection before proceeding.*

This mode-selection line takes precedence over the always-on Full Governance directive in the body of the instruction. Claude will pause before every response and confirm the output mode.

> Verification: Start a new conversation. Claude should immediately ask "Output mode? 1. Full Governance 2. Answer Only" before producing any response.

Platform 3: Gemini (Google)

> **Geographic Restriction**
> The global instructions feature (Instructions for Gemini) is not available in the European Economic Area, Switzerland, or the United Kingdom. Practitioners in those regions should use Gems or the Single-Session Prompt.

Path A — Gem (recommended): Navigate to Explore Gems → New Gem. Name it "HAIA-RECCLIN Governance." Paste the full RECCLIN instruction into the instructions field and save. To use a governed session, open this Gem from the sidebar. The Gem only governs sessions launched through it — standard Gemini conversations are not affected.

Path B — Global instructions (account-wide, where available): Navigate to Settings and help → Personal Intelligence → Instructions for Gemini. Paste the full RECCLIN instruction there.

> Verification: Launch a session through your Gem and ask "What is 2+2?" All ten fields should appear. Confirm you launched via the Gem, not a standard chat window.

Platform 4: Grok (xAI)

Grok's persistent instruction mechanism operates through its memory system. The most reliable method is a direct memory command at the start of a session.

How to load: Open a new Grok conversation and send this message:

> *Remember that I use HAIA-RECCLIN governance in all our sessions. Always respond in Full Governance format with these ten fields in order: Role, Task, Output, Sources,*

Conflicts, Confidence, Expiry, Fact→Tactic→KPI, Recommendation, Decision. Preserve dissent. You suggest. I decide. Human authority is absolute.

Grok will store this instruction in persistent memory and apply it across future sessions. Grok demonstrated strong RECCLIN memory activation in Case Study 006, activating all ten governance fields from stored memory without explicit prompting.

> Verification: Ask "What is 2+2?" All ten fields should appear. If they do not, resend the memory command and confirm Grok acknowledged it before proceeding.

Platform 5: Perplexity

Perplexity supports persistent governance through two mechanisms: AI Profile for account-wide preferences, and Spaces for project-level custom instructions with knowledge sources.

Path A — AI Profile (account-wide): Navigate to your profile icon → Settings → AI Profile. Paste the full RECCLIN instruction there.

Path B — Spaces (recommended for Dispatch and CAIPR workflows): Create a new Space, open its settings, and paste the full RECCLIN instruction into the Custom AI Instructions field. Spaces allow instructions to be stored alongside uploaded documents and knowledge sources.

Perplexity supplemental addition (recommended given its source-validation strength):

For every source cited, provide: summary of the source, the fact or data from that source, the tactic or strategy from that source, and the outcome or KPI from that source. Every source must carry a complete Fact→Tactic→KPI chain to be usable.

This addition activates the original Perplexity Dispatch method that initiated multi-AI governance in 2023. It is an enhancement to the Sources field, not an additional governance field.

> Verification: Ask "What is 2+2?" For AI Profile, confirm the profile was saved before testing. For Spaces, confirm you are working within the Space, not a standard conversation.

Platform 6: Mistral (Le Chat)

Mistral's Le Chat implements persistent instructions through **Agents** — named AI configurations that carry custom instruction sets. For API access, the instruction goes in the system prompt field.

How to load — Le Chat: Navigate to Agents in the left sidebar → Create Agent. Name the agent "HAIA-RECCLIN Governance." Paste the full RECCLIN instruction into the Agent's Instructions field and save. Launch conversations through this Agent.

How to load — API: Place the full RECCLIN instruction in the system prompt field of your API call.

Mistral supplemental addition (recommended given demonstrated tendency toward compressed outputs):

> *Do not summarize or compress your output. Produce full-depth responses in every field. Brevity in the Output field is not a virtue if it sacrifices governance completeness.*

Verification: Launch a conversation through the Agent and ask "What is 2+2?" All ten fields should appear. Confirm you launched via the Agent, not standard Le Chat.

A Note on Memory Reliability

Loading these instructions is the starting point, not the guarantee. Case Study 006 documented that fewer than half of the platforms in the pool activated full governance output from stored memory when the session prompt did not explicitly invoke RECCLIN. For high-stakes sessions, use the Single-Session Prompt as message one even on platforms where custom instructions are already loaded. The mode confirmation costs nothing and verifies that governance is active before output is produced.

For CAIPR sessions, always include the full governance instruction in the dispatch prompt. Do not rely on memory activation for parallel multi-platform review. Memory activation cannot be assumed across all platforms simultaneously.

Maintenance Schedule

UI navigation paths were verified against platform documentation in March 2026. Platform interfaces change frequently. This appendix requires review every quarter, and immediately following any major platform release that affects settings, personalization, agents, instruction mechanisms, or project features. The governance instruction text and field structure are stable. The navigation paths are not.

Single-Session Prompt: For Any Platform

For platforms without persistent instruction support, for practitioners not yet ready for full custom instruction setup, or for any CAIPR dispatch session where memory activation cannot be assumed, paste this prompt as message one. Send it alone. Wait for the mode confirmation. Then send your actual task as message two.

RECCLIN GOVERNANCE MODE — SESSION PROMPT

```
You are operating under HAIA-RECCLIN governance for this session. HAIA-
RECCLIN is a human-AI governance framework developed by Basil C. Puglisi,
MPA. Your role is to produce structured, accountable output that makes your
reasoning visible and verifiable.

Every response you produce must include these ten fields in order. The
initial mode confirmation below is a single-line handshake — ten-field output
begins after mode is confirmed.

Role: Self-assign from: Researcher, Editor, Coder, Calculator, Liaison,
Ideator, Navigator — choose the function that best matches this task.
Task: Repeat my request back in clean language confirming you understood it
correctly.
Output: Your substantive response.
Sources: Citations with sufficient detail to verify. Flag unverified sources
as PROVISIONAL.
Conflicts: Dissent in sources or disagreement with prior outputs. If none
found, state "No conflicts identified" — do not skip this field.
Confidence: 0 to 100 percent with written justification based on evidence
quality.
Expiry: How long this output remains valid. State whether information is
stable or time-sensitive.
Fact→Tactic→KPI: Fact (evidence-grounded finding) then Tactic (executable
action) then KPI (measurable outcome).
Recommendation: Your primary recommendation plus at least one alternative.
Separate your synthesis from the evidence.
Decision: The specific choice requiring my human arbitration. Frame clearly:
accept or challenge, option A or B, or D1/D2/D3 for competing
recommendations. You frame the decision. I make it.

Governance principles: You suggest. I decide. Preserve dissent. Never force
consensus. Human authority is absolute.

If I say "Answer Only," provide a direct response without the governance
structure. Request "reissue in Full Governance format" at any time.

Confirm you are in RECCLIN governance mode by asking: "Output mode? 1. Full
Governance  2. Answer Only" — then wait for my selection.
```

References and Related Documents

- Puglisi, B. C. (2012). Digital Factics: Twitter. Digital Media Press. magcloud.com/browse/issue/471388
- Puglisi, B. C. (2025). Governing AI: When Capability Exceeds Control. ISBN 9798349677687. Amazon.
- Puglisi, B. C. (2025). The Human Enhancement Quotient (HEQ): Measuring Collaborative Intelligence for Enterprise AI Adoption. White Paper v4.3.3. basilpuglisi.com/HEQ
- Puglisi, B. C. (2026a). Measuring Augmented Intelligence: Theoretical Foundations and Empirical Development of the Human Enhancement Quotient (HEQ) and Augmented Intelligence Score (AIS). Working Paper v2.5. basilpuglisi.com/measuring-augmented-intelligence
- Puglisi, B. C. (2026b). Checkpoint-Based Governance: A Constitution for Human-AI Collaboration. Published. basilpuglisi.com
- Puglisi, B. C. (2026c). HAIA-CAIPR Specification v1.1. basilpuglisi.com
- Puglisi, B. C. (2026d). GOPEL: The Code Behind the Policy. Proof of Concept v3.1. basilpuglisi.com / github.com/basilpuglisi/HAIA
- Puglisi, B. C. (2026e). AI Provider Plurality Congressional Package. Documents 1 through 5. basilpuglisi.com
- Puglisi, B. C. (2026f). HAIA-RECCLIN Case Study: The Kimi Outlier. December 2025. basilpuglisi.com
- Puglisi, B. C. (2026g). HAIA Framework Architecture Map v1.8. March 2026. basilpuglisi.com / github.com/basilpuglisi/HAIA
- Puglisi, B. C. (2026h). HAIA-CORE: The Missing Piece in Content Evaluation. basilpuglisi.com
- Puglisi, B. C. (2026i). HAIA-SMART v1.5 Calibration. basilpuglisi.com
- Puglisi, B. C. (2026j). What We Learned: HAIA Multi-AI Practice. basilpuglisi.com

Basil C. Puglisi, MPA

Human-AI Collaboration Strategist

me@basilpuglisi.com | basilpuglisi.com | github.com/basilpuglisi/HAIA

Third Edition | March 2026

© 2026 Basil C. Puglisi. All rights reserved. Human Enhancement Quotient, HEQ, Augmented Intelligence Score, AIS, HAIA, HAIA-RECCLIN, HAIA-CAIPR, GOPEL, Factics, and CBG are trademarks of Basil C. Puglisi. Any use in research, publications, training materials, or commercial applications requires proper attribution. Commercial use requires written authorization.

#AIassisted under HAIA-RECCLIN governance. Human arbiter review and approval required before any external distribution, publication, or application.