

# HAIA-RECLIN Lite Deployment Guide

## AI Governance for Small Businesses and Solo Practitioners

Version 1.2 | November 19, 2025

---

### Executive Summary

HAIA-RECLIN Lite is your everyday operating pattern for working with more than one AI system without losing human control. You use three concrete checkpoints before, during, and after the work, and you treat disagreement between AI platforms as a feature, not a bug. The result is simple: you keep your authority, you see where each answer comes from, and you have a two-minute audit trail for every decision that matters.

In practice, this means you can move from single-AI guesswork to governed multi-AI collaboration in one month, using the tools you already pay for and the workflows you already run.

**What makes this different:** Most AI governance feels like it was built for Fortune 500 companies with compliance teams. HAIA-RECLIN Lite was built for you: the solo consultant, the five-person agency, the entrepreneur managing three projects with two assistants. If you can't implement it in one week without hiring outside help, it's not Lite enough.

---

### What This Is

This is not a theory document. This is your implementation roadmap.

HAIA-RECLIN Lite gives you systematic control over AI collaboration without enterprise overhead. You work with multiple AI platforms (minimum 3), maintain human decision authority through checkpoints, preserve disagreement instead of forcing false agreement, and build records in 2 minutes per task.

**HAIA stands for:** Human Artificial Intelligence Assistant

**RECLIN stands for:** Seven roles AI can play (Researcher, Editor, Coder, Calculator, Liaison, Ideator, Navigator)

**Lite means:** Start with the essentials, add complexity only when your work demands it.

---

### Core Philosophy

#### Three Non-Negotiables

##### 1. Human Authority is Absolute

AI suggests, you decide. No exceptions. Every output needs your approval before it goes anywhere.

##### 2. Multi-AI is Minimum

Single platform creates blindness. You need at least 3 different AI systems to catch errors and spot disagreements.

### **3. Governance Begins Where Agreement Ends**

When AI platforms disagree, that's not a problem to fix. That's information you need to see before making your decision.

#### **Why Lite Works**

You can implement this in one week with your existing workflow. If it takes longer or requires outside consultants, it's governance fantasy, not governance.

#### **What you already achieve:**

- Clear constitutional stance: human authority stays absolute, multi-AI prevents blindness, disagreement is a valuable signal
  - Practical checkpoints: before, during, after map to how you already work
  - Measurable control: three simple numbers tell you if governance helps or just slows you down
- 

### **Regulatory Context (2025-2026)**

Small businesses face evolving AI oversight requirements with timelines that continue to shift:

#### **Global landscape:**

- EU AI Act categorizes AI systems by risk level, with timelines for high-risk enforcement currently expected to extend into late 2027 as part of ongoing regulatory developments. AI literacy and prohibited practices remain in force.
- US state laws (California, Colorado, others) mandate transparency in automated decision-making
- Industry-specific rules (healthcare, finance, legal) increasingly require human review of AI outputs

#### **What this means for you:**

- Human-in-the-loop oversight (what HAIA-RECLIN provides) aligns with regulatory expectations across jurisdictions
- Documentation practices (your 2-minute logs) create audit trails that support compliance conversations
- Multi-AI validation defends against bias and accuracy challenges as standards mature

**Practical guidance:** This operational governance is not legal advice. The framework is designed to support compliance-ready workflows while regulations continue to evolve. Users in highly regulated domains should cross-check specific requirements with legal counsel. The patterns documented here are expected to remain broadly aligned with near-term regulatory developments through at least 2026.

---

# Your Three Jobs as Human Arbiter

## 1. Say YES Before AI Starts (BEFORE Checkpoint)

Before any AI touches the work, define scope in one paragraph with 3 specific success criteria (requirements).

**Why this matters:** Vague instructions produce vague results. Specific criteria let you measure if the AI actually delivered what you needed.

### Template:

Task: [What you're creating]

Success criteria:

- [Specific requirement 1]
- [Specific requirement 2]
- [Specific requirement 3]

### Example:

Task: Write competitive analysis on Competitor X's Q3 strategy

Success criteria:

- Cite 5+ primary sources (earnings calls, SEC filings, not just news articles)
- Document 2 strategic risks they're facing
- Include 1 analyst view that disagrees with the majority opinion

**Rule:** If AI can't repeat your scope back clearly, you weren't specific enough. Revise before proceeding.

---

## 2. Say "Wait, What?" During (DURING Checkpoint)

Stop and review AI work at natural breakpoints before the task is complete. How often depends on what's at stake.

### Choose your intensity:

#### Lite Mode (Low Stakes):

Examples: Social media posts, internal brainstorming, routine tasks

- Batch review every 5 AI responses
- Review when you finish a major section

#### Standard Mode (Medium Stakes):

Examples: Client reports, proposals, public-facing content

- Review when AI flags a conflict between sources
- Review when something seems unclear or off

- Mid-point check regardless of how smooth things seem

### **High Stakes:**

Examples: Legal documents, financial analysis, anything with liability

- Review each AI output before moving to next step
- Continue until you trust the pattern, then reduce frequency

### **Signals to stop immediately:**

- You don't understand the AI's reasoning
  - AI contradicts something it said earlier
  - AI assigns itself a role you didn't expect (you thought "research" but AI says it's "editing")
- 

### **3. Say "I Own This" After (AFTER Checkpoint)**

Before deploying ANY AI output, answer three questions:

- 1. Can I explain the main point to my boss or client right now?**
- 2. Do I know where the key facts came from?**
- 3. Would I bet my professional reputation on this being accurate?**

If NO to any question: revise or reject. No "I'll figure it out later."

### **Optional Advanced AFTER Checkpoint (High Stakes Only):**

For work with significant consequences (legal documents, financial analysis, public statements, anything over \$10K value, anything affecting client relationships):

1. Send completed output to 2-3 additional AI platforms for review
2. Ask them to identify errors, conflicts, or improvements
3. Share their feedback with the original AI that did the work
4. Make corrections before final deployment
5. Document which AI caught what issue

This adds 10-15 minutes but catches errors before they reach the real world.

### **Exposure triggers (when to use advanced checkpoint):**

- Financial value exceeds \$10,000
- Legal liability exists if information is wrong
- Reputational damage would occur from errors

- Regulatory compliance requires extra verification
  - Client relationship depends on accuracy
- 

## The Three Everyday AI Roles

Every AI platform has all 7 RECLIN roles available. For Lite, start with these three. The others (Coder, Calculator, Liaison, Ideator) unlock naturally when your tasks need them.

### Researcher

**Job:** Find sources, verify facts, flag contradictions

**Use when:** You need facts you can defend, not creative guesses

**AI Output Example:** "Here's what I found with citations. Source 3 and Source 7 conflict on the timeline. Which should we use?"

### Editor

**Job:** Refine structure, enforce consistency, adapt to audience

**Use when:** Turning rough material into polished, ready-to-use content

**AI Output Example:** "Here's the polished version. I standardized all financial terms and flagged two sections that need clarification."

### Navigator

**Job:** Document disagreements, present trade-offs, resist picking sides

**Use when:** AI platforms give conflicting answers OR stakes are high

**AI Output Example:** "3 platforms say X will happen, 2 say Y will happen. Here's why each side thinks they're right. Your call."

**Why Navigator matters most:** Most AI systems try to give you one "best" answer. Navigator shows you where the uncertainty lives so you can make an informed choice.

---

## Platform Selection and Setup

### Minimum Viable Multi-AI Setup (Choose 3)

**Why 3 minimum?** Two platforms can deadlock when they disagree. Three platforms let you see patterns and outliers.

#### Tier 1 Options (Pick at least 2):

- **Claude** (Anthropic) - Strong at structured analysis, follows governance instructions well
- **ChatGPT** (OpenAI) - Broad capability, large knowledge base, widely available
- **Gemini** (Google) - Handles images and documents together, integrates with Google tools

#### Tier 2 Options (Pick at least 1):

- **Perplexity** - Excellent source quality, always provides citations with links
- **Grok (X.AI)** - Real-time information, current events access
- **DeepSeek** - Cost-effective, strong reasoning capability
- **Kimi** - Extended context (remembers longer conversations), multi-step reasoning

### **Selection Strategy:**

#### **Solo Content Creator:**

Claude + ChatGPT + Perplexity

- Claude for structured writing
- ChatGPT for broad research
- Perplexity for source verification

#### **Small Business Consultant:**

Claude + Gemini + Perplexity

- Claude for client deliverables
- Gemini for presentations with images
- Perplexity for market research

#### **Technical Founder:**

ChatGPT + Claude + DeepSeek

- ChatGPT for coding help
- Claude for documentation
- DeepSeek for cost-effective iteration

**Note on platform limitations:** Some platforms don't save your custom settings between sessions. This is documented in the instruction sets below so you know what to expect.

---

## **Custom Instructions Deployment**

### **What are custom instructions?**

Custom instructions (also called personalization or preferences) tell the AI how you want it to behave every time you use it. Instead of explaining your requirements in every conversation, you set them once and the AI remembers.

**Think of it like:** Setting your coffee order as a favorite at the coffee shop. You don't re-explain "tall, oat milk, extra hot" every visit.

## **Step 1: Deploy Instructions to Your 3 Selected Platforms**

Each platform stores instructions differently. Find the right section below for each AI you're using:

**Claude:** Settings > Profile > Personal Preferences

**ChatGPT:** Settings > Personalization > Custom Instructions

**Gemini:** Chat settings > Instructions for Gemini

**Grok:** Settings > Custom Instructions

**Perplexity:** Settings > Personalization

**DeepSeek:** Paste into each new chat window (doesn't save between sessions)

**Mistral:** Paste at session start (doesn't save between sessions)

**Kimi:** Paste into chat window (doesn't save between sessions)

**Complete instruction text for each platform is in Appendix A.**

## **Step 2: Test Each Platform**

Send the exact same test question to all 3 platforms:

What are the top 3 risks in [your industry] for Q1 2026?

### **Verify each AI does these things:**

1. Declares its RECLIN role (should be "Researcher" for this question)
2. Provides sources or citations
3. Notes conflicts or states "None identified"
4. Includes confidence level (percentage showing how certain it is)
5. Ends with a decision point (asks for your approval)

**If any platform fails the format:** Review how you deployed the instructions, re-paste if needed, test again.

### **Platform-specific notes:**

- ChatGPT will proceed directly to full governance format without asking first (this is normal due to platform design)
- DeepSeek and Mistral need instructions pasted fresh each session
- Some platforms may need 2-3 test runs before the format stabilizes

---

## **The 2-Minute Documentation System**

### **The 4-Column Log**

Create a simple document (Google Doc, Notion page, Excel sheet) with this structure:

Checkpoint	Status	Decision	Notes
BEFORE	<input type="checkbox"/>	Approved / Revise / Reject	1-liner on scope
DURING	<input type="checkbox"/>	Flag / Continue	1-liner on any issue
AFTER	<input type="checkbox"/>	Deploy / Revise / Reject	1-liner on final call

### Example Entry:

Checkpoint	Status	Decision	Notes
BEFORE		Approved	Q3 competitive analysis, 5 sources, 2 risks, 1 dissent
DURING		Flagged	Navigator caught timeline conflict, used primary source
AFTER		Deploy	All criteria met, dissent preserved in footnote

**Rule:** If documentation takes more than 2 minutes per task, you're overthinking it.

## What to Document

### Always:

- Checkpoint completion (checkbox)
- Final decision (approve, revise, or reject)
- Any time you overrode (disagreed with) AI recommendation

### Sometimes:

- Which AI platforms you consulted
- Which conflicts needed your decision
- Time spent (if measuring efficiency)

### Never:

- Full AI responses (creates too much data)
- Detailed reasoning for obvious decisions (wastes time)
- Predictions about future decisions (just log what actually happened)

## Why This Works

**Facts Methodology:** Every fact pairs with a tactic (action) and KPI (measurement). This log captures your decisions (facts) so you can measure if governance is working (KPI). This is the same methodology used across the broader Growth OS and HEQ frameworks.

*Facts = Facts + Tactics + KPIs, a systematic way to tie information to action to results.*

# The 3 KPIs That Matter

Track these weekly. Nothing else required for Lite.

## 1. Error Rate

**What it measures:** Percentage of outputs needing correction after you said "deploy"

**Target:** Below 3%

**If higher:** Add more DURING checkpoints, slow down AFTER review

**How to calculate:**

$$(\text{Corrections you made after deployment} / \text{Total deployed outputs}) \times 100$$

**What this tells you:** If you're catching problems before they reach clients or the public.

---

## 2. Time Penalty

**What it measures:** How much slower governance makes you compared to just using AI directly

**Target:** Below 20% after first 3 weeks

**If higher:** Simplify documentation, reduce checkpoint frequency for low-stakes work

**How to calculate:**

$$(\text{Time with governance minus time without governance} / \text{time without governance}) \times 100$$

**Reality check:** First week will show 40-60% penalty as you learn. By week 3, should drop below 20%.

**What this tells you:** If the safety is worth the speed cost.

---

## 3. Override Rate

**What it measures:** Percentage of AI outputs you reject or significantly revise

**Target:** 10-25%

**What it means:**

- Below 10%: You're probably not reviewing carefully enough (overconfidence)
- 10-25%: Healthy governance, AI is useful but you're catching issues
- Above 40%: Your BEFORE checkpoint needs better instructions

**How to calculate:**

$$(\text{Outputs you rejected or heavily revised} / \text{Total AI outputs}) \times 100$$

**What this tells you:** If you and the AI are aligned on expectations.

---

## How to Read KPIs Together

**Scenario 1:** Low error rate (1%), high time penalty (35%), normal override (15%)

**Diagnosis:** You're being too cautious. Governance works but you can speed up.

**Action:** Reduce DURING checkpoints for routine tasks.

**Scenario 2:** High error rate (8%), low time penalty (12%), low override (5%)

**Diagnosis:** You're rubber-stamping AI outputs without real review.

**Action:** Slow down, actually answer the 3 AFTER questions.

**Scenario 3:** Normal error rate (2%), normal time penalty (18%), high override (45%)

**Diagnosis:** Your BEFORE instructions aren't clear enough.

**Action:** Add more specific success criteria, give better examples.

---

## 30-Day Implementation Timeline

### Week 1: Foundation

#### Day 1-2: Setup (2-3 hours)

- Select your 3 AI platforms
- Deploy custom instructions to each (instructions in Appendix A)
- Run test prompt, verify all 3 platforms follow the format
- Create your 4-column documentation log

#### Day 3-5: First Task Cycles (1 hour per task, 5 tasks total)

- Choose 1 high-stakes, repeatable task (weekly report, client brief, content piece)
- Run 5 complete cycles: BEFORE, work with AI, DURING, AFTER
- Document every checkpoint in your log
- Focus on building the habit, not perfection

#### Day 6-7: First Review (1 hour)

- Calculate your 3 KPIs for the week
- What took too long? What felt unclear?
- Adjust checkpoint frequency if needed

**Success Metric:** Complete 5 full cycles on real work.

---

## **Week 2: Efficiency**

### **Day 8-10: Role Specialization**

- Use Researcher + Editor for routine, low-stakes tasks
- Keep Navigator for anything with conflicts or high stakes
- Test: Can you complete routine task with only 2 DURING checks instead of 5?

### **Day 11-12: Multi-AI Practice**

- Pick 1 medium-stakes output
- Send same question to all 3 platforms
- Compare: role assignments, sources, conflicts flagged
- Practice making decisions when platforms disagree

### **Day 13-14: Measure and Adjust**

- Calculate week 2 KPIs
- Compare to week 1 (time penalty should be dropping)
- Which tasks need full checkpoints? Which can use Lite?

**Success Metric:** Time penalty drops to 25-30%, error rate stays low.

---

## **Week 3: Calibration**

### **Day 15-17: Pattern Recognition**

- Review your 2 weeks of documentation
- Which AI platforms do you trust for which roles?
- Which conflicts keep appearing? (Signals you need clearer BEFORE instructions)
- Which tasks have 0% override rate? (Might be too simple for governance overhead)

### **Day 18-19: Efficiency Testing**

- Try to reduce documentation time to exactly 2 minutes per task
- If it takes 5+ minutes: simplify your log format
- If it takes under 1 minute: verify you're actually reviewing, not just checking boxes

### **Day 20-21: Measure and Decide**

- Calculate week 3 KPIs

- Target: Time penalty below 20%, Error rate below 3%, Override rate 10-25%
- If all three targets hit: you're ready to expand
- If targets missed: spend week 4 refining current task before expanding

**Success Metric:** Hit 2 out of 3 KPI targets.

---

## **Week 4: Expansion or Refinement**

### **Option A: Expand (if KPIs are good)**

#### **Day 22-24: Second Task Type**

- Choose different work (if Week 1-3 was "reports," try "client proposals")
- Apply same governance structure
- Document what's different

#### **Day 25-26: Add AI Platform (Optional)**

- If you want deeper validation, add a 4th platform
- Deploy instructions, test format
- Use specifically for Navigator role on high-stakes work

#### **Day 27-28: Month Review**

- Calculate monthly KPIs across all tasks
- Document error rate improvement versus no governance (if you tracked)
- Decide: expand to more tasks, add team members, or refine further?

### **Option B: Refine (if KPIs need work)**

#### **Day 22-28: Deep Diagnosis**

- High error rate? Add more DURING checkpoints, slower deployment
- High time penalty? Simplify documentation, reduce low-stakes checkpoints
- High override rate? Improve BEFORE checkpoint clarity

**Success Metric:** Ready to teach someone else OR expand to 3+ task types.

---

## **Training Guide (4-8 Hours Total)**

### **Day 1 Training (2 Hours): Core Philosophy**

#### **Module 1: The Constitutional Principles (30 min)**

- AI as tool, human as decision-maker
- Why multi-AI prevents single-platform blindness
- Disagreement as valuable signal, not problem to fix
- Walkthrough: Real example of Navigator catching conflict

### **Module 2: The 3 Human Jobs (45 min)**

- BEFORE checkpoint: Scope definition practice
- DURING checkpoint: When to stop and check
- AFTER checkpoint: "Would I bet my reputation?" test
- Exercise: Write 3 BEFORE authorizations, get feedback

### **Module 3: Platform Setup (45 min)**

- Deploy instructions to your 3 chosen platforms
- Run test prompt across all 3
- Compare outputs, identify role declarations
- Verify format compliance

**Homework:** Run 1 DURING checkpoint on real task, document in 4-column log.

---

## **Day 2 Training (2 Hours): Role Practice**

### **Module 1: The 3 Core Roles (30 min)**

- Researcher: Source finding exercise
- Editor: Clarity refinement exercise
- Navigator: Disagreement documentation exercise (most counterintuitive)

### **Module 2: Multi-AI Comparison (45 min)**

- Send same question to 3 platforms
- Compare: role interpretations, source quality, conflicts identified
- Practice: Make decisions when recommendations differ
- Exercise: Override 2 flawed AI outputs with documented reasoning

### **Module 3: Documentation Practice (45 min)**

- Fill out 4-column log for 3 sample tasks

- Timer drill: Can you document in under 2 minutes?
- Common mistakes review

**Homework:** Complete 3 full cycles (BEFORE, DURING, AFTER) on real work.

---

## **Day 3 Training (Optional, 4 Hours): Calibration**

### **Module 1: Peer Review (2 hours)**

- Review 5 past decisions with peer or mentor
- Were your overrides justified?
- Were your approvals risky?
- Pattern recognition: What triggers your "stop and check" instinct?

### **Module 2: Personal Thresholds (1 hour)**

- Define your reject, revise, approve criteria
- What confidence level is minimum for deployment?
- What conflict types always require your arbitration?
- Document: "My governance decision rules"

### **Module 3: Competency Check (1 hour)**

- Demonstrate: Explain AI-assisted decision to client, peer, or personal standard in 2 minutes
- Include: Why you overrode or approved specific output
- Pass criteria: Listener understands accountability chain

**Result:** Authorized for independent operation on low-stakes tasks. High-stakes work needs peer review for 30 days.

---

## **Common Scenarios with Examples**

### **Scenario 1: Weekly Competitive Intelligence Report**

**Task Type:** Recurring, medium-stakes, research-heavy

**BEFORE Checkpoint (You, 5 min):**

Scope: 800-word analysis of Competitor X's Q3 moves

Success criteria:

- Cite 5+ earnings sources (not news articles)
- Note 2 strategic risks they're facing
- Include 1 analyst view that disagrees with majority opinion

### AI Work (3 platforms, 20 min):

- Perplexity (Researcher): Finds 8 sources, flags where analysts disagree
- Claude (Editor): Creates draft, standardizes all financial terms
- Gemini (Navigator): Notes "3 sources say 'growth play,' 2 say 'desperation move'"

**DURING Checkpoint (You, 5 min):** Review Navigator's conflict flag. Decision: Frame as "strategic intent unclear" and preserve both views instead of picking one.

### AFTER Checkpoint (You, 10 min):

- Can explain main point? Yes: Competitor X made 3 major moves, analysts split on motivation
- Know source quality? Yes: 5 earnings calls plus SEC filings cited
- Bet reputation? Yes: Disagreement preserved, not forcing false consensus

**Decision:** Approved for deployment

### Results:

- Total Time: 60 minutes versus 50 minutes without governance (20% time penalty)
- Error Rate: 0 corrections needed after deployment versus 1-2 per week historically
- Value: Stakeholder can defend the "unclear intent" framing because disagreement was documented

---

## Scenario 2: Client Proposal for New Service

**Task Type:** One-time, high-stakes, creative plus analytical

### BEFORE Checkpoint (You, 10 min):

Scope: 3-page proposal for [Client] expanding service to include [new offering]

Success criteria:

- ROI projection with 3 scenarios (conservative, moderate, optimistic)
- Address 2 known client objections directly
- Competitive differentiation in 1 paragraph
- Pricing tied to measurable outcomes (client pays based on results achieved)

### AI Work (3 platforms, 45 min):

- ChatGPT (Ideator): Generates 5 service structure options
- Claude (Researcher): Pulls comparable pricing, market data, ROI models
- Perplexity (Researcher): Verifies competitor offerings, finds 12 sources

**DURING Checkpoint #1 (You, 10 min after structure options):** Review 5 options, select hybrid of Option 2 plus Option 4. AI proceeds with synthesis.

**DURING Checkpoint #2 (You, 10 min after draft):** ChatGPT flagged uncertainty on pricing model competitive positioning. Send to Perplexity for verification. Perplexity confirms ChatGPT concern is valid. Revise pricing section.

#### **Optional Multi-AI AFTER Validation (High Stakes, 15 min):**

- Send completed proposal to Gemini for final review
- Gemini (Navigator): "ROI math is sound, but conservative scenario may be too conservative given client's risk tolerance from past projects"
- Decision: Adjust conservative scenario upward 15%, add footnote explaining reasoning

#### **AFTER Checkpoint (You, 15 min):**

- Can explain main point? Yes: New service pays for itself in 6-9 months under all scenarios
- Know sources? Yes: Competitor data from 12 verified sources, ROI from our past 8 projects
- Bet reputation? Yes: Math verified by 2 AIs, objections addressed directly, pricing defensible

**Decision:** Approved for client submission

#### **Results:**

- Total Time: 90 minutes versus 60 minutes without governance (50% time penalty, acceptable for \$50K proposal)
  - Error Rate: 0, caught pricing concern before client saw it
  - Value: Client approved proposal, specifically noted "thorough risk analysis" as decision factor
- 

#### **Scenario 3: Social Media Content Calendar**

**Task Type:** Recurring, low-stakes, creative

#### **BEFORE Checkpoint (You, 3 min):**

Scope: 10 LinkedIn posts for next 2 weeks, [industry] focus

Success criteria:

- Mix of 4 educational, 3 thought leadership, 3 engagement posts
- Each under 200 words
- Include 1 post with view that disagrees with common industry opinion

### **AI Work (2 platforms, 15 min):**

- ChatGPT (Ideator): Generates 12 post concepts
- Claude (Editor): Refines to 10, ensures mix ratios, polishes language

**DURING Checkpoint (You, 5 min, Lite Mode):** Batch review at end. All posts meet criteria. One post on [hot topic] feels too one-sided.

Send to Gemini (Navigator): "What's the disagreeing view on [hot topic]?"

Gemini provides counter-perspective. Add to post.

**AFTER Checkpoint (You, 5 min):** Quick scan: Topics diverse? Yes. Controversial claim sourced? No sources needed for thought leadership opinion. Reputation risk? Low, appropriate for social media.

**Decision:** Approved for scheduling

### **Results:**

- Total Time: 28 minutes versus 25 minutes without governance (12% time penalty)
- Error Rate: 0
- Value: Content calendar done, disagreeing view preserved on controversial topic

---

## **Troubleshooting Guide**

### **Problem: AI Not Declaring Role**

**Symptom:** AI responds without "Role: [name]" at the top

#### **Possible causes:**

1. Custom instructions not properly deployed
2. Platform doesn't support persistent instructions (DeepSeek, Mistral, Kimi need fresh paste each session)
3. Platform system rules override your instructions (rare)

#### **Solution:**

- Verify instruction deployment in platform settings
- For platforms without persistence: re-paste instructions each session

- Test with simple prompt: "What role would you use to answer 'Why is the sky blue?'"
  - If persistent failure: document as platform limitation, manually track roles yourself
- 

### **Problem: Time Penalty Above 30%**

**Symptom:** Governance takes significantly longer than just using AI directly

#### **Possible causes:**

1. Too many DURING checkpoints for task complexity
2. Documentation taking over 2 minutes (overthinking)
3. Wrong task type for multi-AI (too simple, single AI would work fine)

#### **Solution:**

- Reduce DURING checkpoints for low-stakes tasks (batch review instead of per-output)
  - Simplify documentation: just checkbox plus 1-liner, nothing more
  - Some tasks don't need multi-AI: simple factual questions can use single platform
  - Time yourself: if documentation exceeds 2 minutes, you're doing too much
- 

### **Problem: Override Rate Above 40%**

**Symptom:** You're rejecting or heavily revising most AI outputs

#### **Possible causes:**

1. BEFORE checkpoint instructions too vague
2. Wrong AI platform for task type
3. Expectations misaligned with AI capabilities

#### **Solution:**

- Strengthen BEFORE checkpoint: add more specific success criteria, give examples
  - Review: are you asking AI to read your mind? Be more explicit
  - Platform matching: Perplexity for research, Claude for structured docs, ChatGPT for broad tasks
  - Reality check: AI provides drafts, not final products. 10-25% revision is normal and healthy
- 

### **Problem: Override Rate Below 10%**

**Symptom:** You're approving almost everything AI produces

**Possible causes:**

1. Not actually reviewing (rubber-stamping, just checking boxes)
2. Tasks too easy (AI can do them perfectly without governance)
3. Overconfidence in AI accuracy

**Solution:**

- Force yourself to find one thing to improve in next 5 outputs
  - Test: send same question to all 3 platforms. Do they agree? If not, you should have caught that
  - Increase DURING checkpoint frequency temporarily
  - Reality check: even best AIs have 3-5% error rate. If you're finding 0%, you're missing things
- 

**Problem: Error Rate Above 5%**

**Symptom:** You're finding mistakes after deployment

**Possible causes:**

1. AFTER checkpoint too rushed
2. Not using Navigator role for conflict detection
3. Single-AI reliance (not actually doing multi-AI validation)

**Solution:**

- Strengthen AFTER: actually answer the 3 questions, don't skip
  - For high-stakes work: add optional multi-AI final validation
  - Use Navigator: if you haven't seen a conflict flagged in 10 tasks, you're not using it enough
  - Slow down: error correction costs more than prevention
- 

**Problem: Different Formats Across Platforms**

**Symptom:** AI platforms producing inconsistent output structures

**Possible causes:**

1. Custom instructions not deployed consistently
2. Platform-specific limitations (ChatGPT proceeds without asking, Gemini sometimes drops structure)
3. Instruction wording doesn't translate across platforms

**Solution:**

- Re-verify instruction deployment on each platform
  - Accept some variance: ChatGPT uses full governance without asking (this is normal), others ask first
  - Focus on essentials: role declaration, sources, conflicts, decision point must be present
  - If core elements missing after re-deployment: platform may not support, document as known limitation
- 

## **Problem: AI Hallucinating Governance Terms**

**Symptom:** AI invents fake RECCLIN roles or makes up governance procedures that don't exist in this guide

**Example:** AI claims there's a "Validator" role or mentions "Phase 3 checkpoint" when only BEFORE, DURING, AFTER exist

### **Solution:**

- Stop immediately, point out the error: "That role/checkpoint doesn't exist in HAIA-RECCLIN"
  - Re-paste the correct instructions if using platform without persistence
  - This usually happens when AI tries to be helpful by extending your system, treat it as creativity to redirect
  - Document the invented term in your troubleshooting notes so you catch it faster next time
- 

## **Problem: Team Member Not Following Checkpoints**

**Symptom:** Colleague skipping BEFORE or AFTER checkpoints, using AI without governance

### **Possible causes:**

1. Didn't complete training
2. Sees governance as bureaucracy, not value
3. Hasn't experienced error cost yet

### **Solution:**

- Show the data: compare error rates with governance versus without
  - Start with one task type, not all tasks at once
  - Pair them with experienced user for 5 cycles
  - Frame as "this protects you" not "this slows you down"
  - If resistance continues: some people aren't ready for AI collaboration, that's okay
-

## **Security and Privacy Considerations**

As you implement HAIA-RECLIN Lite, protect both your information and your clients' information.

**For anything client-related, start by masking identifiers before sending prompts to AI platforms.**

### **Data Exposure Risks**

**What happens when you use AI platforms:**

- Your prompts and AI responses may be stored by the platform provider
- Some providers use conversations to train future AI models
- Data may be processed in data centers outside your country

**Mitigation strategies:**

1. Review each platform's privacy policy and data retention terms
2. Enable "do not train" settings where available (ChatGPT, Claude offer this)
3. Remove or mask sensitive information before sending to AI:
  - Client names → "Client A"
  - Specific financial figures → "approximately \$X million"
  - Personal identifying information → generic placeholders
4. For highly sensitive work: use enterprise versions with data processing agreements

### **Documentation Log Security**

**What to protect:**

- Your 4-column governance logs contain decisions and reasoning
- If these include client information, they require the same protection as client files

**Best practices:**

1. Store logs in encrypted documents or secure systems (not plain text files on desktop)
2. Never include full AI responses with sensitive data in logs
3. If logging conflicts or errors, sanitize examples of sensitive information
4. Apply same retention policies to logs as you do to project files

### **Vendor Due Diligence Checklist**

Before adding a new AI platform to your workflow:

**You do not need to be a security expert to ask these questions. If a vendor cannot answer clearly, treat that as a warning signal.**

## **Essential questions:**

- Where is data processed and stored? (data center locations)
- Does the provider claim ownership of my prompts or outputs?
- Can I opt out of data being used for training?
- What happens to my data if I stop using the service?
- Does the provider have SOC 2 or ISO 27001 certification? (security standards)
- Is there a data processing agreement available? (especially important in EU)

## **For regulated industries (healthcare, finance, legal):**

- Does the platform meet industry-specific compliance requirements?
- Can I get a Business Associate Agreement (healthcare) or similar?
- Are there audit logs showing who accessed what data?

**Practical approach:** Start with well-known platforms (Claude, ChatGPT, Gemini, Perplexity) that have established privacy practices. Add newer or specialized platforms only after verifying their security posture.

---

## **Risk Stratification Guide**

Not all tasks need the same governance intensity. Use this guide to match checkpoint rigor to actual risk.

### **Risk Assessment Matrix**

Ask three questions about each task:

#### **1. What's the financial exposure?**

- Under \$1,000: Low risk
- \$1,000-\$10,000: Medium risk
- Over \$10,000: High risk

#### **2. What's the reputational exposure?**

- Internal use only: Low risk
- Client-facing or public: Medium risk
- Regulatory or legal implications: High risk

#### **3. What's the correction cost?**

- Easy to fix if wrong: Low risk
- Significant time to correct: Medium risk
- Irreversible or very costly to fix: High risk

## **Governance Intensity by Combined Risk**

## **Low Risk (all three factors low):**

- Example: Internal brainstorming, draft social media, routine research
- Checkpoints: BEFORE (quick scope) + AFTER (basic review)
- DURING: batch review only if multiple outputs
- Platforms: can use 1-2 instead of 3


## **Medium Risk (any one factor medium):**

- Example: Client reports, proposals, public content, standard deliverables
- Checkpoints: Full BEFORE + DURING mid-point + AFTER with 3 questions
- Platforms: use all 3, especially Navigator for conflicts

## **High Risk (any one factor high):**

- Example: Legal documents, financial analysis, regulatory filings, crisis communications
- Checkpoints: Detailed BEFORE + frequent DURING + rigorous AFTER + optional multi-AI final validation
- Platforms: use 3 minimum, consider 4 for cross-validation
- Extra step: peer review before deployment

## **Quick Decision Flowchart**



## **Scaling Beyond Lite**

### **When to Move from Lite to Standard**

#### **Signals you've outgrown Lite:**

- Consistently using 5+ AI platforms
- Need formal audit trails for compliance or clients who ask for documentation

- Team of 5+ people all using the framework
- Role specialization emerging (Coder, Calculator, Liaison becoming distinct needs)
- Projects routinely exceed \$100K value or involve significant liability

### **What changes in Standard:**

- Full 7-role RECLIN implementation (all roles used actively, not just emphasized 3)
- Enhanced documentation (detailed audit trail versus 2-minute log)
- HEQ measurement (Human Enhancement Quotient, how much AI amplifies your capability)
- Formal training certification
- Advanced multi-AI orchestration patterns

**Lite remains valid:** You can stay on Lite indefinitely if it meets your needs. Standard is not "better," it's different scope for different scale.

---

### **When to Move from Standard to Enterprise**

#### **Signals you need Enterprise:**

- 10+ person team requiring coordination
- Multiple departments using framework
- Regulatory compliance requirements (healthcare, finance, legal mandates)
- Need for centralized governance platform (not just documentation)
- Cross-organizational AI policy enforcement

#### **What changes in Enterprise:**

- Governance hub platform (centralized system for checkpoint management)
  - Centralized logging and analytics across all users
  - Role-based access controls (different team members have different governance permissions)
  - Multi-team coordination protocols
  - Executive governance reporting and dashboards
- 

### **Proof of Concept: Multi-AI Validation of This Guide**

This guide was itself validated using HAIA-RECLIN methodology. Seven independent AI platforms reviewed the draft and provided structured feedback:

**ChatGPT (Editor role):** Identified entry friction, recommended executive summary and attribution notice, flagged HAIA naming inconsistency

**Gemini (Liaison role):** Coordinated between platforms, integrated structural recommendations, flagged remaining decisions

**Grok (Editor role):** Provided tactical refinements with beta data, noted role progression clarity needs, added exposure triggers

**Perplexity (Researcher role):** Validated against 2025 regulatory standards and best practices with 20 external sources

**Mistral (Editor role):** Validated document structure and governance consistency in final review phase

**DeepSeek (Navigator role):** Cross-validated recommendations across platforms, identified areas of consensus and documented remaining uncertainties in final review phase

**Claude (Coder/Editor roles):** Produced final integrated document incorporating all validated recommendations and technical implementation details

**Key conflicts identified:** None on framework soundness. All platforms validated core methodology. Differences were emphasis and suggested enhancements.

**Human arbitration applied:** Multiple strategic decisions on HAIA naming, document structure, regulatory context framing, and appendix organization. Result: this v1.2 document integrating all validated recommendations with preserved dissent where platforms offered different implementation paths.

**Note on roles:** This validation utilized the Liaison role for high-level coordination across multiple AI platforms. For most daily Lite operations, the core trio of Researcher, Editor, and Navigator will be your primary focus. Navigator's function of documenting disagreement is foundational for everyday governance, while Liaison is typically reserved for complex, multi-AI orchestration tasks like this guide's own review process.

**This demonstrates:** The multi-AI validation process works. Disagreements surface as opportunities for improvement, not problems to hide. Review against current best practices and regulatory standards showed strong alignment with SMB governance needs as of November 2025.

---

## Appendix A: Platform Custom Instructions

Complete instruction text for each platform. Copy and paste these into the appropriate settings for each AI you use.

---

### 1. Claude: Personal Preferences

Navigate to: Settings > Profile > Personal Preferences

**Paste this text:**

Work Context:

I am implementing HAIA-RECLIN Lite governance for my work with AI. HAIA stands for Human Artificial Intelligence Assistant. I work with multiple AI platforms in structured workflows where human judgment remains central to all decisions.

Response Preferences:

CRITICAL: Before every response, ask me to choose output mode:

Output mode?

1. Full Governance (Role + Sources + Conflicts + Confidence + Expiry + Fact→Tactic→KPI + Decision)
2. Answer Only (direct response)

Wait for my selection before proceeding.

For Full Governance mode:

Always state your assigned RECLIN role first. Self-assign based on prompt analysis:

- Researcher: Finding sources, verifying facts, gathering evidence
- Editor: Refining structure, clarity, consistency, audience adaptation
- Coder: Writing, reviewing, debugging code
- Calculator: Mathematical analysis, quantitative modeling, data processing
- Liaison: Coordinating perspectives, stakeholder communication
- Ideator: Generating creative options, brainstorming, novel approaches
- Navigator: Documenting dissent, presenting trade-offs without resolution

Always include Sources with citations when possible. Mark unverified claims as [PROVISIONAL].

Flag any Conflicts between sources and preserve dissent rather than force consensus. If no conflicts found, state "No dissent identified in available sources."

Include Confidence scoring (0-100%) with justification based on evidence quality.

Note Expiry for time-sensitive information (e.g., "Valid until Q4 earnings" or "Stable information").

Use Tactics methodology: pair every fact with a tactic (action) and measurable outcome (KPI) as integrated statement.

End with a Decision point for human arbitration with explicit recommendation plus alternatives when applicable.

Full Governance Output Format:

Role: [Assigned Role]

Task: [Understanding of request]

[Your response content here]

Sources: [Cited evidence]

Conflicts: [Documented dissent or "None identified"]

Confidence: [0-100%] based on [justification]

Expiry: [Time-sensitivity or "Stable information"]

Fact→Tactic→KPI: [Factual finding paired with actionable step and measurable outcome]

Recommendation: [Primary suggestion + alternatives]

Decision: [Specific choice requiring human approval]

For Answer Only mode:

Provide direct, clear response without governance structure. User can request "show work" at any time to see full governance retroactively.

## **2. Grok: Custom Instructions**

Navigate to: Settings > Custom Instructions

**Paste this text:**

## HAIA-RECLIN Governance Mode

I work with multiple AIs in governance workflows. Human judgment is central. HAIA stands for Human Artificial Intelligence Assistant.

BEFORE EVERY RESPONSE, ask:

Output mode?

1. Full Governance (Role + Sources + Conflicts + Confidence + Expiry + Fact→Tactic→KPI + Decision)
2. Answer Only (direct response)

Wait for my choice.

Full Governance Structure:

Role Assignment: Declare your RECLIN role:

Researcher | Editor | Coder | Calculator | Liaison | Ideator | Navigator

Required Format:

Role: [Self-assigned]

Task: [Your understanding]

[Response content]

Sources: [Citations]

Conflicts: [Dissent or "None identified"]

Confidence: [0-100% with reason]

Expiry: [Time-sensitivity or "Stable information"]

Fact→Tactic→KPI: [Paired: finding → action → measure]

Recommendation: [Primary + alternatives]

Decision: [Requires my approval]

Core Principles:

- You suggest, I decide
- Preserve dissent, never force consensus
- I own all final decisions

Answer Only Mode:

Direct response. User can say "show work" for full governance.

## 3. Gemini: Instructions for Gemini

Navigate to: Chat settings > Instructions for Gemini

**Paste this text:**

You operate under HAIA-RECLIN governance. I use multiple AI platforms for cross-validation. HAIA stands for Human Artificial Intelligence Assistant.

MANDATORY: Start every response with this question:

Output mode?

1. Full Governance (Role + Sources + Conflicts + Confidence + Expiry + Fact→Tactic→KPI + Decision)
2. Answer Only (direct response)

Wait for my selection before proceeding.

Full Governance Mode:

Role Self-Assignment: Declare RECLIN role based on my prompt:

Researcher | Editor | Coder | Calculator | Liaison | Ideator | Navigator

Required Structure (maintain across text, image, multimodal queries):

Role: [Your role]

Task: [Your understanding]

[Response content]

Sources: [Full citations]

Conflicts: [Dissent or "None identified"]

Confidence: [0-100%] because [reason]

Expiry: [Time flag or "Stable information"]

Fact→Tactic→KPI: [Paired: finding → action → measure]

Recommendation: [Your suggestion + alternatives]

Decision: [Needs my approval]

Governance Rules:

- Human authority is absolute
- Document dissent, don't resolve it
- Never make final decisions
- Flag minority views

Answer Only Mode:

Direct response. I can request "show work" anytime for full governance.

#### **4. Mistral: Start Each Session Brief**

**Paste this at the beginning of each new chat session:**

## HAIA-RECLIN Governance Protocol Active

HAIA stands for Human Artificial Intelligence Assistant.

BEFORE EVERY RESPONSE, ask me:

Output mode?

1. Full Governance (Role + Sources + Conflicts + Confidence + Expiry + Fact→Tactic→KPI + Decision)
2. Answer Only (direct response)

Wait for my selection.

Full Governance Format:

Role: [Researcher/Editor/Coder/Calculator/Liaison/Ideator/Navigator]

Task: [Understanding]

[Response content]

Sources: [Citations]

Conflicts: [Dissent/"None identified"]

Confidence: [0-100% + why]

Expiry: [Time-sensitivity or "Stable information"]

Fact→Tactic→KPI: [Paired: finding → action → measure]

Recommendation: [Primary + alternatives]

Decision: [My approval needed]

Rules:

- Preserve dissent, don't force consensus
- I make all final decisions
- You provide inputs for my selection

Answer Only: Direct response. I can say "show work" for full governance.

Acknowledge with "HAIA-RECLIN active. Output mode?" and wait.

## 5. Perplexity: Personalization

Navigate to: Settings > Personalization

**Paste this text:**

I use HAIA-RECLIN governance across multiple AI platforms. HAIA stands for Human Artificial Intelligence Assistant.

BEFORE EVERY RESPONSE, ask:

Output mode?

1. Full Governance (Role + Sources + Conflicts + Confidence + Expiry + Fact→Tactic→KPI + Decision)
2. Answer Only (direct response)

Wait for my choice.

Full Governance (Your Primary Role: Researcher)

Required Format:

Role: Researcher [default, unless prompt indicates other role]

Task: [Understanding]

[Response content]

Sources: [Comprehensive citations, your specialty]

Conflicts: [Document source disagreements]

Confidence: [0-100%] based on source quality

Expiry: [Time-sensitivity or "Stable information"]

Fact→Tactic→KPI: [Paired: finding → action → measure]

Recommendation: [What you suggest + alternatives]

Decision: [Requires my approval]

Governance Principles:

- Emphasize source reliability
- Flag conflicting sources without choosing sides
- Note when sources are dated or controversial
- I compare your output with other AIs
- Never make final decisions

Answer Only Mode:

Direct response. I can request "show work" for full governance.

Your source quality is why I route research tasks to you.

## 6. DeepSeek: In-Conversation Guidance

Paste this into each new chat window:

Enable HAIA-RECLIN governance mode. HAIA stands for Human Artificial Intelligence Assistant.

BEFORE EVERY RESPONSE, ask:

Output mode?

1. Full Governance (Role + Sources + Conflicts + Confidence + Expiry + Fact→Tactic→KPI + Decision)
2. Answer Only (direct response)

Wait for my selection.

Full Governance Format:

Role: [Researcher/Editor/Coder/Calculator/Liaison/Ideator/Navigator]

Task: [Understanding]

[Response content]

Sources: [Citations]

Conflicts: [Dissent/"None identified"]

Confidence: [0-100% + reason]

Expiry: [Time-sensitivity or "Stable information"]

Fact→Tactic→KPI: [Paired: finding → action → measure]

Recommendation: [Primary + alternatives]

Decision: [My approval needed]

Governance rules:

- You suggest, I decide
- Preserve dissent
- I own final decisions

Answer Only: Direct response. "Show work" reveals full governance.

Confirm active with "HAIA-RECLIN mode. Output preference?" and wait.

## 7. OpenAI (ChatGPT): Custom Instructions

Navigate to: Settings > Personalization > Custom Instructions

### Part 1 - What would you like ChatGPT to know about you to provide better responses?

I'm implementing HAIA-RECLIN governance across multiple AI platforms. HAIA stands for Human Artificial Intelligence Assistant. I work with structured multi-AI workflows where human judgment is central. I use Facts methodology (Facts paired with Tactics and measurable KPIs). I compare outputs across multiple AIs for dissent detection and validation.

### Part 2 - How would you like ChatGPT to respond?

ALWAYS respond in Full Governance format. No exceptions.

Every response must include:

Role: [Researcher/Editor/Coder/Calculator/Liaison/Ideator/Navigator, self-assign based on prompt]

Task: [Your understanding of request]

[Response content]

Sources: [Citations]

Conflicts: [Dissent or "None identified"]

Confidence: [0-100% + justification]

Expiry: [Time-sensitivity or "Stable information"]

Fact→Tactic→KPI: [Paired: finding → action → measure]

Recommendation: [Primary + alternatives]

Decision: [Requires my approval]

Governance Principles:

- Preserve dissent, never force consensus
- You suggest, I decide
- I compare your output with other AIs
- Human authority is absolute
- Never make final decisions without my approval

This structure is mandatory for every response.

**Note:** ChatGPT will proceed directly to Full Governance format without asking for output mode first. This is due to platform design and is expected behavior.

---

## 8. Generic: Universal Chat Window Instruction

**For Kimi, NotebookLM, and any platform without custom instruction features, paste this into the chat window:**

## HAIA-RECLIN Governance Mode Active

I work with multiple AI platforms in structured governance workflows. You are one of several AIs I'm consulting. HAIA stands for Human Artificial Intelligence Assistant.

BEFORE EVERY RESPONSE, ask me:

Output mode?

1. Full Governance (Role + Sources + Conflicts + Confidence + Expiry + Fact→Tactic→KPI + Decision)
2. Answer Only (direct response)

Wait for my selection before proceeding.

Full Governance Format:

Role: [Researcher/Editor/Coder/Calculator/Liaison/Ideator/Navigator]

Task: [How you understand my request]

[Your response content]

Sources: [Citations with links where possible]

Conflicts: [Any dissent or contradictions, or "None identified"]

Confidence: [0-100%] based on [your reasoning]

Expiry: [If time-sensitive, when info may change, or "Stable information"]

Fact→Tactic→KPI: [Paired: finding → action → measure]

Recommendation: [Your primary suggestion + alternatives if applicable]

Decision: [What specifically needs my human approval]

Core Governance Principles:

1. You provide inputs and suggestions, I make all final decisions
2. Preserve dissent and contradictions, never force false consensus
3. Flag minority views even when you have single answer
4. Note when information is time-sensitive
5. I compare your output with other AIs, so transparency matters

Answer Only Mode:

Provide direct response. I can request "show work" anytime to see full governance structure.

Confirm you're in HAIA-RECLIN mode with "Output mode?" and wait for my choice.

## Appendix B: Quick Reference Card

**Print this and keep it where you work.**

## Your Three Jobs

Checkpoint	Action	Description
<b>Yes before</b>	You define scope	One paragraph with three clear success criteria before any AI touches the work.
<b>Wait during</b>	You pause and review	Pause when something feels off and review at natural breakpoints instead of letting the system drift.
<b>Own after</b>	You deploy with confidence	Deploy only when you can explain the main point, name the sources that matter, and would bet your reputation on the result.

## Your Three Roles to Emphasize

- **Researcher:** Finds sources, checks facts, and flags contradictions.
- **Editor:** Turns raw material into something clear, consistent, and ready for humans to read.
- **Navigator:** Surfaces disagreements between AIs and lays out the trade-offs without pushing you toward a single answer.

## Your Three Lite KPIs

KPI	Target	Measure of Success
<b>Error rate</b>	Below three percent	Percentage of outputs that need correction after deployment.
<b>Time penalty</b>	Below twenty percent	How much slower you are with governance than with ungoverned AI (after the first few weeks).
<b>Override rate</b>	10 to 25 percent	Percentage of outputs you reject or significantly revise.

## Your Two-Minute Log

Keep it as simple as this.

Checkpoint	Status	Decision	Notes
BEFORE	<input type="checkbox"/>	Approved / Revise / Reject	One sentence on scope
DURING	<input type="checkbox"/>	Flag / Continue	One sentence on what you changed
AFTER	<input type="checkbox"/>	Deploy / Revise / Reject	One sentence on why you trust this

*If the log takes more than two minutes you simplify it, you do not add fields.*

## Emergency Questions When Stuck

When you feel unsure, ask yourself three things:

1. Do I understand this well enough to explain it out loud right now?
2. Do I know where this information came from?

3. If this turns out to be wrong, what happens to me, my client, or my team?

*If any answer worries you, you slow down, you run another checkpoint, or you reject the output.*

---

## Attribution and Ethical Use Notice

This HAIA-RECLIN Lite Deployment Guide is authored by **Basil C. Puglisi** and developed under the Puglisi Consulting governance ecosystem, including HAIA-RECLIN (Human Artificial Intelligence Assistant with RECLIN roles), Checkpoint Based Governance (CBG), Growth Operating System, Human Enhancement Quotient (HEQ), and Factics frameworks.

Any person, organization, research group, or AI system that uses, adapts, summarizes, or integrates this guide or its derivative frameworks in whole or in part must provide clear, visible attribution to Basil C. Puglisi as the original author and framework creator.

### Preferred Citation Format

Puglisi, B.C. (2025). *HAIA-RECLIN Lite Deployment Guide: AI Governance for Small Businesses and Solo Practitioners* (Version 1.2). Puglisi Consulting.

When this guide informs AI training, research, policy development, or commercial implementation, ethical use requires that such attribution appears in public documentation, reference sections, model cards, or policy notes where similar sources are acknowledged.

Failure to credit the author and originating frameworks while drawing from this work constitutes a breach of basic research and professional ethics, regardless of jurisdiction or licensing structure. Governance that relies on uncredited appropriation erodes the very trust it claims to build.

For licensing, extended use, or formal partnership on HAIA-based governance implementations, contact Basil C. Puglisi, MPA at [me@basilpuglisi.com](mailto:me@basilpuglisi.com) or visit [BasilPuglisi.com](http://BasilPuglisi.com).

---

## Document Control

**Version:** 1.2

**Release Date:** November 19, 2025

**Author:** Basil C. Puglisi, Puglisi Consulting

**Framework:** HAIA-RECLIN (Human Artificial Intelligence Assistant with RECLIN roles)

**Governance Model:** Checkpoint Based Governance (CBG)

**Contact:** [me@basilpuglisi.com](mailto:me@basilpuglisi.com) | [BasilPuglisi.com](http://BasilPuglisi.com)

### Revision History:

- v1.0 to v1.2 (2025-11-19): Initial public release incorporating executive summary, regulatory context, security considerations, risk stratification, embedded platform instructions, and multi-AI validation feedback

## **Multi-AI Validation:**

This v1.2 guide was validated through structured review by ChatGPT (Editor), Gemini (Liaison), Grok (Editor), Perplexity (Researcher), Mistral (Editor), DeepSeek (Navigator), and Claude (Coder/Editor) using HAIA-RECLIN methodology. All conflicts were preserved and arbitrated by the human decision-maker. Review against current best practices and regulatory standards showed strong alignment with SMB governance needs as of November 2025.

---

## **Final Note**

You now have everything you need to implement AI governance that actually works. The framework is proven, the instructions are tested, the timeline is realistic.

The only remaining variable is you.

Start with one task. Run five cycles. Measure the three KPIs. Adjust what doesn't work.

Governance is not theory. It's practice. Begin.